

TECHNOLOGY

Workloads



A Day in the Life of an DevOps Engineer

You are working as a DevOps Engineer in an organization, and you have been asked to design a solution to create a highly available MySQL and WordPress deployment using autoscaling. All the users should be added using ConfigMaps and the sensitive data should be added using Secrets.

You also need to ensure that the service is running on the NodePort and the WordPress Pod should not deploy if the MySQL service is not deployed.

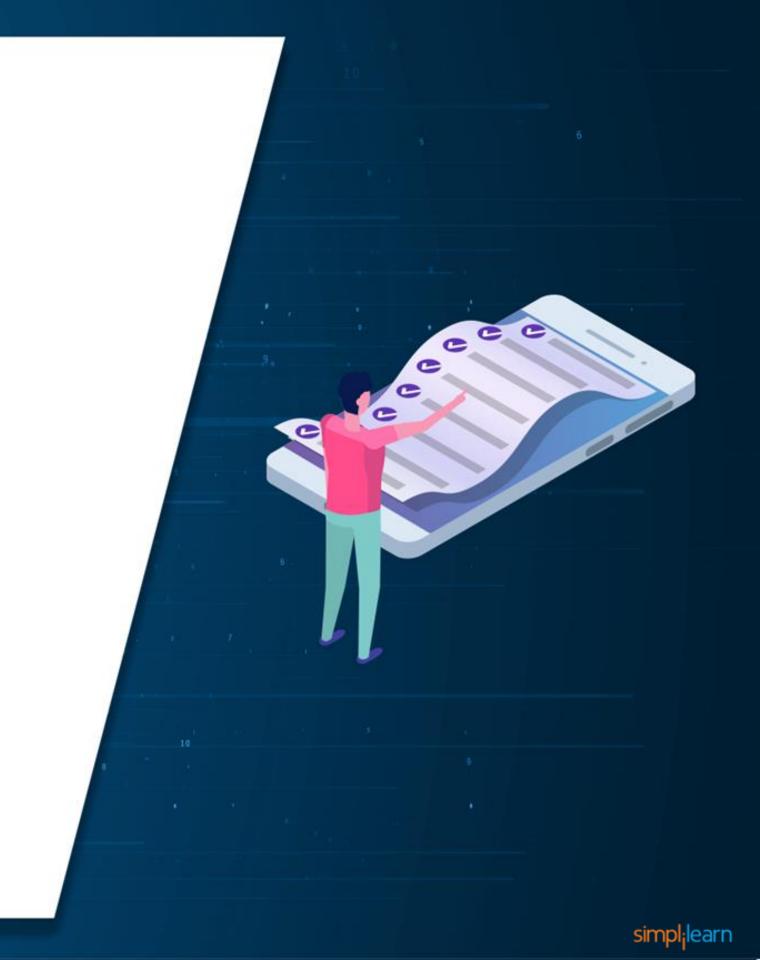
To achieve all the above, along with some additional concepts, we would be learning a few concepts in this lesson that will help you find a solution for the above scenario.



Learning Objectives

By the end of this lesson, you will be able to:

- Deploy a multitier application
- Create a Pod using an init container
- Work on Pod allocation
- Develop a horizontal Pod autoscaler
- Create a Kubernetes ReplicaSet



TECHNOLOGY

Overview of Workloads

Overview

A workload is an application running on Kubernetes.



Whether the workload has a single or several components that work together, it runs inside a set of Pods.



When a Pod is running in a cluster, a critical fault on the node fails all Pods on that node.

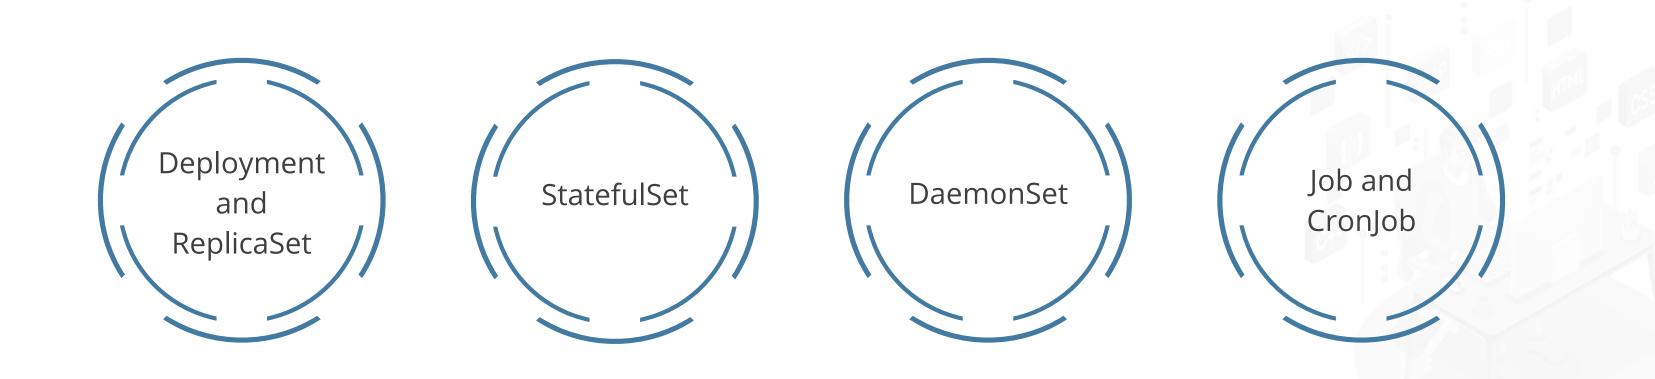


Kubernetes treats that level of failure as final. It creates a new Pod to recover, even if the node becomes healthy later.



Overview

Kubernetes provides several built-in workload resources:



Users can add in a third-party workload resource if they want a specific behavior that is not part of Kubernetes' core.



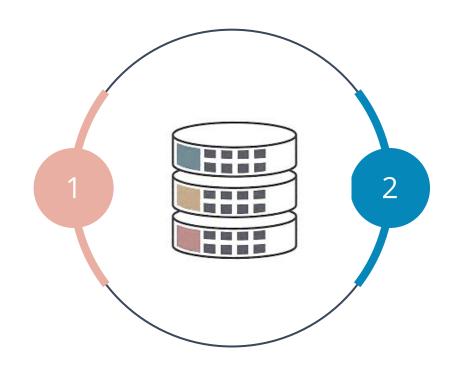
TECHNOLOGY

Deployment

What Is Deployment?

A Deployment provides declarative updates for Pods and ReplicaSets.

Users can describe the desired state in a Deployment, and the Deployment Controller changes the actual state to the desired state at a controlled rate.



Users can define Deployments to create new ReplicaSets or remove existing Deployments and adopt all their resources with new Deployments.

Use Cases

Typical use cases for Deployments are:

- 1 Create a Deployment to rollout a ReplicaSet
- 5 Pause the rollout of a Deployment

2 Declare the new state of Pods

6 Use the status of the Deployment

3 Rollback to an earlier Deployment revision

- 7 Clean up older ReplicaSets
- 4 Scale up the Deployment to facilitate more load



Deploying Multitier Application Using Kubernetes



Duration: 20 mins

Problem Statement:

You've been asked to deploy a multitier application with Kubernetes.

Assisted Practice: Guidelines

Steps to be followed:

- 1. Creating a Deployment for MySQl
- 2. Creating a Deployment for WordPress
- 3. Creating a Service for WordPress and MySQL Deployment
- 4. Verifying the Deployment of the application



Deploying Multitier Application Postgres and Gogs



Duration: 20 mins

Problem Statement:

You've been asked to deploy a multitier application Postgres and Gogs using Kubernetes

oor

Assisted Practice: Guidelines

Steps to be followed:

- 1. Creating a Deployment for Postgres
- 2. Creating a Deployment for Gogs
- 3. Creating a Service for Postgres and Gogs Deployment
- 4. Verifying the Deployment of the application



Deploying a Voting Application



Duration: 20 mins

Problem Statement:

You've been asked to deploy a voting application in Kubernetes.

Assisted Practice: Guidelines

Steps to be followed:

- 1. Creating a namespace
- 2. Creating the application
- 3. Verifying the Deployment of the application



TECHNOLOGY

Understanding Pods



Introduction

Pods are the smallest deployable units of computing that users can create and manage in Kubernetes.



A Pod is a group of one or more containers with shared storage and network resources and a specification for how to run the Containers.



A Pod's contents are always co-located, co-scheduled, and run in a shared context.



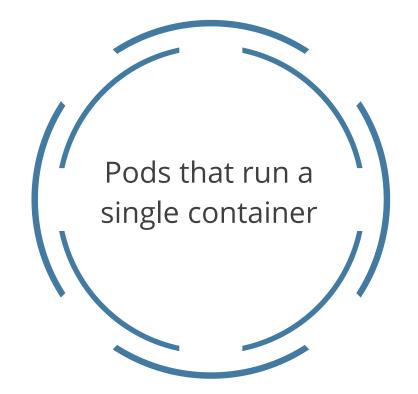
A Pod models an application-specific **logical host**. It contains one or more application containers that are relatively tightly coupled.



Using Pods

Users do not need to create Pods directly, even singleton Pods. Instead, create them using workload resources such as Deployment or Job. If the Pods need to track the state, consider the StatefulSet resource.

Users use Pods in a Kubernetes cluster in two main ways:

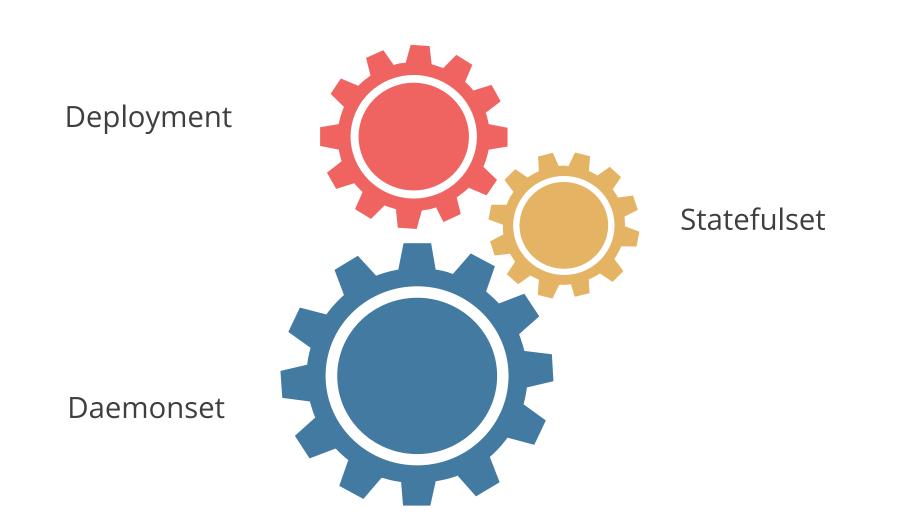


Pods that run multiple containers that need to work together

Pods and Controllers

When a Pod gets created, the new Pod is scheduled to run on a node in the cluster. A controller for the resource handles replication and rollout and automatic healing in case of Pod failure.

Some examples of workload resources that manage one or more Pods:

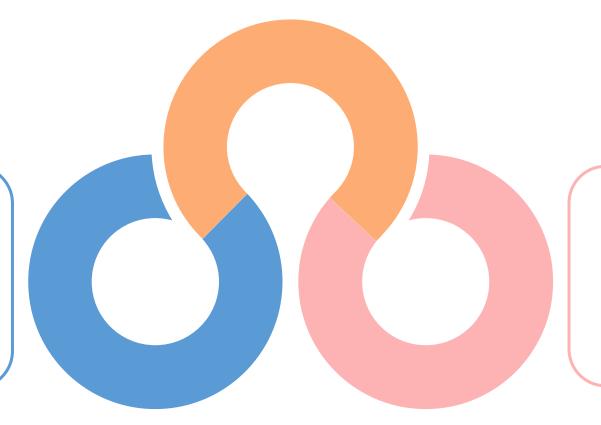




Pods and Controllers

Each workload resource implements its own rules for handling changes to the Pod template.

Modifying the Pod template or switching to a new Pod template has no direct effect on the Pods that already exist.



The Kubelet does not directly observe or manage any of the details around Pod templates and updates.

Pod Templates

Pod Templates are specifications that help users to create Pods. They are included in workload resources such as Deployments, Jobs, and DaemonSets.

The following sample is a manifest for a basic Job with a template that starts a single container:

```
apiVersion: batch/v1
Kind: Job
Metadata:
 name: hello
 spec:
  ktemplate:
   #This is the pod template
   spec:
    containers:
 - name: hello
  protocol: TCP
  image: busybox
  command: ['sh', '-c', 'echo "Hello, kubernetes!" && sleep 3600']
restartPolicy: OnFilure
   # The pod template ends here
```



Pod Update and Replacement

Pod update operations like **patch** and **replace** have some limitations:



The metadata about a Pod is immutable.



If the **metadata.deletionTimestamp** is set, no new entry can be added to the metadata.finalizers list.



Pod updates may not change fields other than spec.containers[*].image, spec.initContainers[*].image, spec.activeDeadlineSeconds, or spec.tolerations.



Only two types of updates are allowed while updating the **spec.activeDeadlineSeconds** field.



Resource Sharing and Communication

Pods enable data sharing and communication among their constituent containers.

Storage in Pods

All containers in the Pod can access the shared volumes, allowing those containers to share data.

Pod networking

Within a Pod, containers share an IP address and port space and can find each other via localhost.



Privileged Mode for Containers

Any container in a Pod can enable Privileged Mode using the privileged flag on the security context of the container spec.



This is useful for containers that want to use operating system administrative capabilities such as manipulating the network stack or accessing hardware devices.

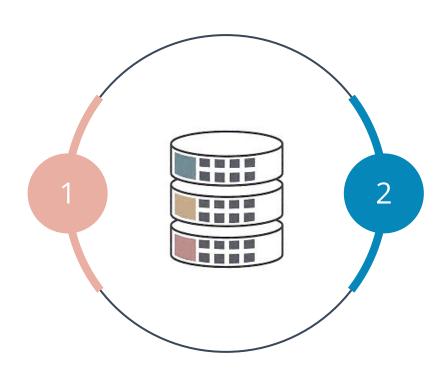
TECHNOLOGY

Pod Life Cycle

Pod Life Cycle

Pods follow a defined lifecycle, starting in the **Pending** phase, moving through the **Running** phase, and then through either the **Succeeded** or **Failed** phase, depending on whether any container in the Pod terminated in failure.

Kubernetes tracks different container statuses within a Pod and decides what action to take to restore the Pod's health.



In the Kubernetes API,
Pods have both a
specification and an
actual status.

Pods are Ephemeral

Pods are created, assigned a unique ID (UID), and scheduled to nodes where they remain until termination (according to restart policy) or deletion.



If a node dies, the Pods scheduled to that node are scheduled for deletion after a timeout period.



Kubernetes uses a higher-level abstraction called controller, which handles the work of managing the relatively disposable Pod instances.



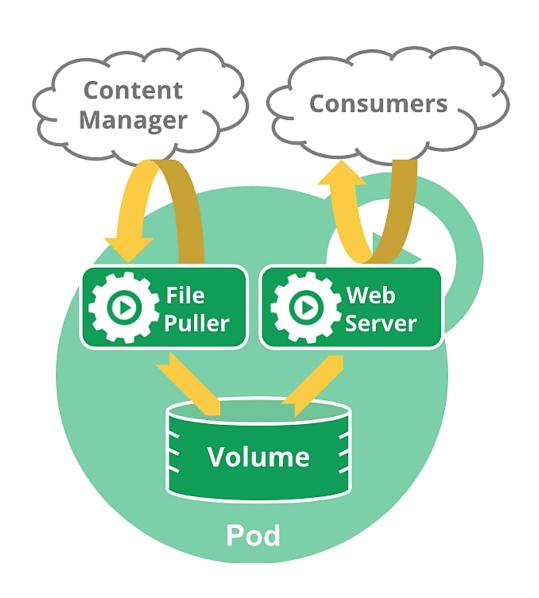
A given Pod (as defined by a UID) is never **rescheduled** to a different Node.



A Pod can be replaced by a new, near-identical Pod with even the same name, if desired but with a different UID.

Pods and Volumes

A multi-container Pod contains a file puller and a web server that use a persistent volume for shared storage between the containers.





Phase of a Pod

The phase of a Pod is a simple, high-level summary of where the Pod is in its life cycle. A Pod's status field is a **PodStatus** object, which has a **phase** field. The following are the possible values for a phase:

Value	Description
Pending	The Pod has been accepted by the Kubernetes cluster, but one or more of the containers has not been set up and made ready to run.
Running	The Pod has been bound to a node, and all the containers have been created.
Succeeded	All containers in the Pod have terminated successfully and will not be restarted.

Pod Phase

If a node dies or is disconnected from the rest of the cluster, Kubernetes applies a policy for setting the phase of all Pods on the lost node to **failed**. The following are the possible values for a phase:

Value	Description
Failed	All containers in the Pod have terminated, and at least one container has terminated in failure.
Unknown	This phase typically occurs due to an error in communicating with the node where the Pod should be running.

Container States

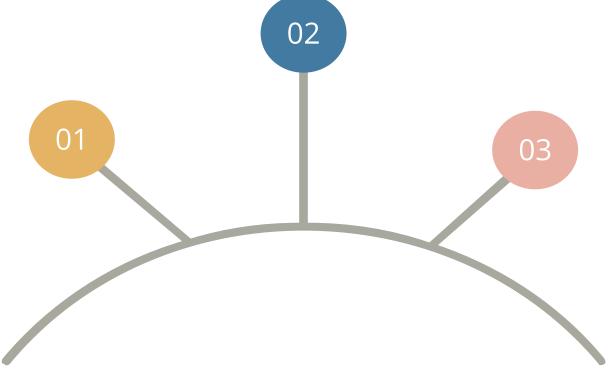
Kubernetes tracks the state of each container inside a Pod. Each state has a specific meaning:

Running

The Running status indicates that a container is executing without issues.

Waiting

If a container is not in either the Running or Terminated state, it is Waiting.

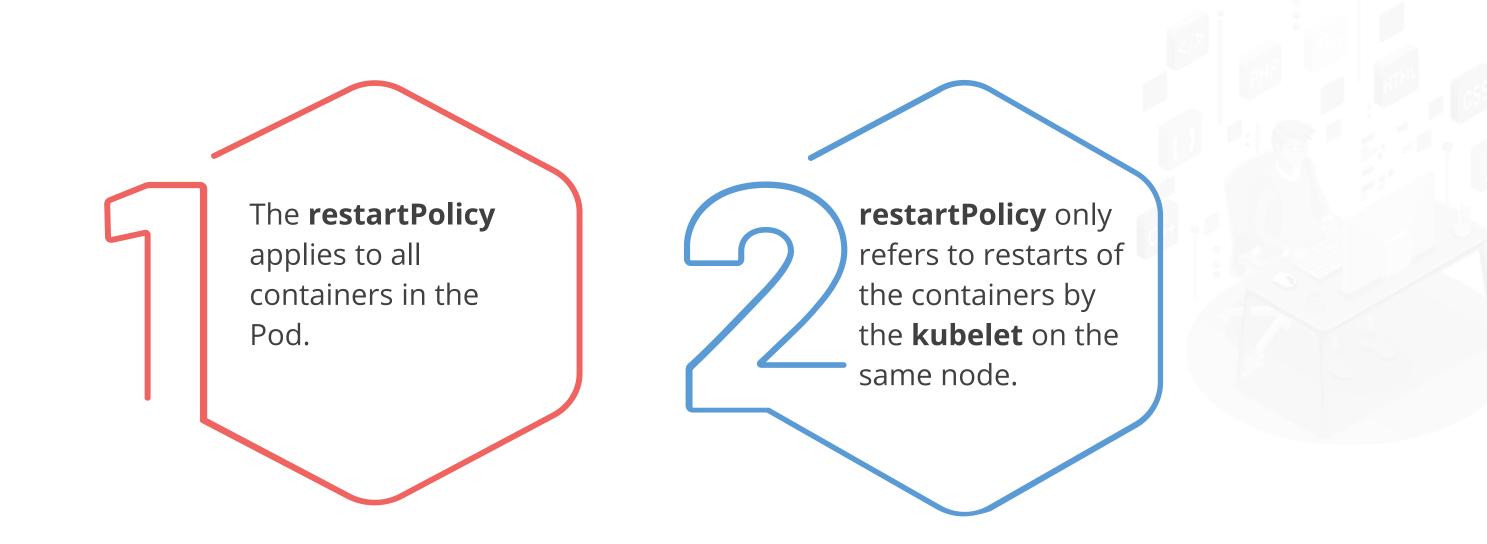


Terminated

A container in the terminated state began execution and then either ran to completion or failed for some reason.

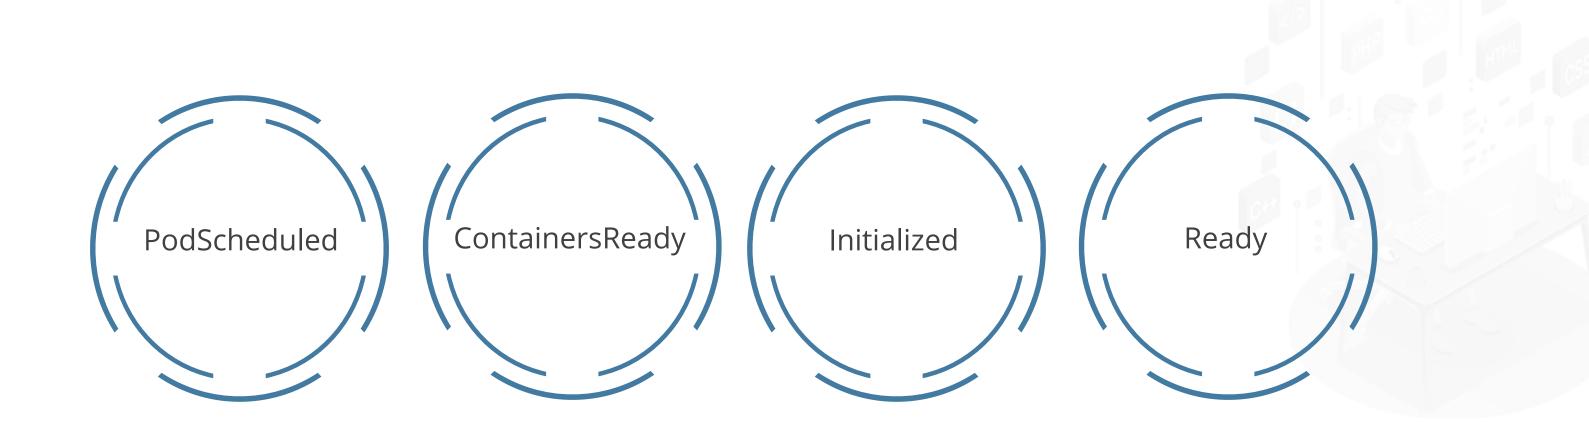
Container Restart Policy

The **spec** of a Pod has a **restartPolicy** field with possible values **Always, OnFailure,** and **Never**.



Pod Conditions

A Pod has a **PodStatus**, which has an array of **PodConditions** through which the Pod has passed or not. Following are the PodConditions:



Pod Conditions

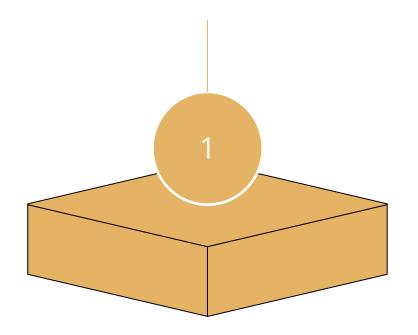
Field Name	Description
type	Represents the name of the Pod condition
status	Indicates whether that condition is applicable, with possible values, such as True , False , or Unknown
lastProbeTime	Represents the timestamp of when the Pod condition was last probed
lastTransitionTime	Represents the timestamp for when the Pod last transitioned from one status to another
reason	Represents machine-readable , UpperCamelCase text indicating the reason for the condition's last transition
message	Represents human-readable message indicating details about the last status transition

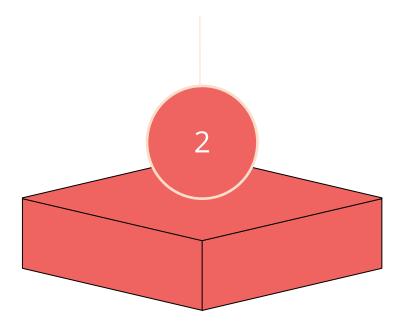
Pod Readiness

An application can inject extra feedback or signals into PodStatus: **Pod readiness**.

To use this, set **readinessGates** in the Pod's **spec** to specify a list of additional conditions that the **kubelet** evaluates for Pod readiness.

Readiness gates are determined by the current state of **status.condition** fields for the Pod.





Pod Readiness

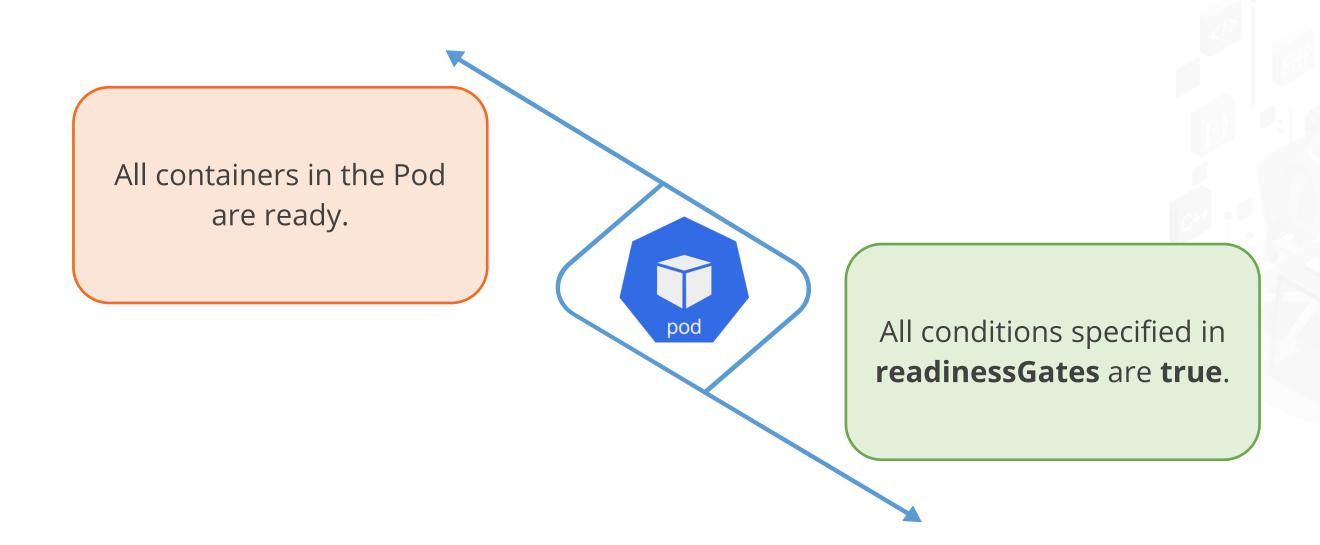
Here is an example:

```
Kind: pod
spec:
 readinessGates:
  - conditionType: "www.example.com/feature-1"
status:
  conditions: /
                       # a built in podCondition
  - type:: Prefix
   status:
    lastProbeTime:
    LastTransitionTime: 2018-01-01T00:00:00Z
    - tyoe: 'www.example.com/feature-1" # an extra podCondition
    status: "False"
    lastProbeTime: null
    lastTransitionTime: 2018-01-01T00:00:00Z
  containerStatuses:
  - containerID: docker://abcd...
    Ready: true
```



Status for Pod Readiness

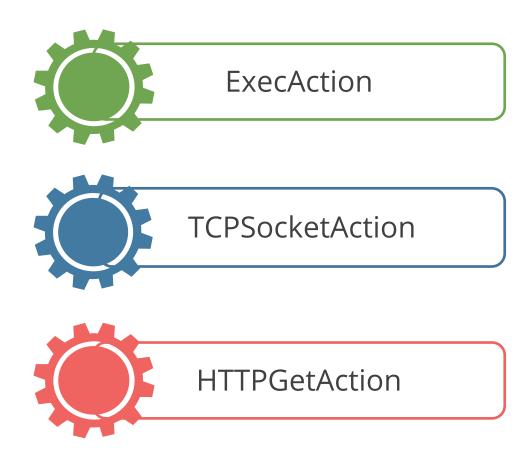
In case of a Pod that uses custom conditions, it is evaluated to be ready only when both the statements given below apply:



Container Probes

A probe is a diagnostic performed periodically by the kubelet on a container.

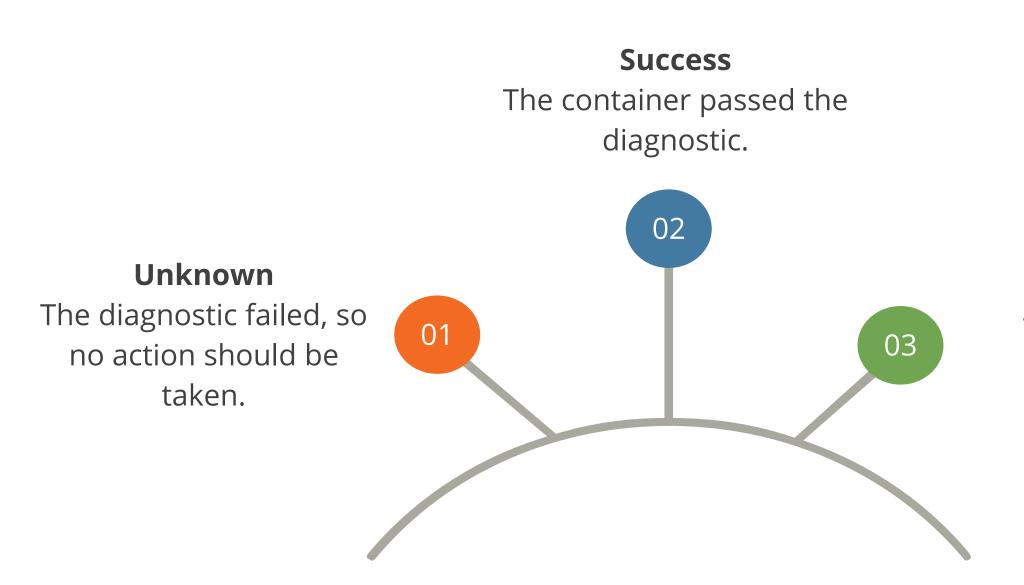
The kubelet can invoke different actions:





Container Probes

Each probe has one of the three results:



Failure

The container failed the diagnostic.

Container Probes

The kubelet can optionally perform and react to three kinds of probes on running containers:

livenessProbe

Indicates whether the container is running

readinessProbe

Indicates whether the container is ready to respond to requests

startupProbe

Indicates whether the application within the container is started



Termination of Pods

The design aim is to be able to request deletion, know when processes terminate, and ensure that deletes are eventually complete.



If the kubelet or the container runtime's management service is restarted while waiting for processes to terminate, the cluster retries from the start, including the full original grace period.



Forced Pod Termination

Forced deletions can be potentially disruptive for some workloads and their Pods.

The control plane cleans up terminated Pods when the number of Pods exceeds the configured threshold.

7

The **kubectl delete** command supports the **--grace- period=<seconds>** option, which allows overriding the default values.

When a forced deletion is performed, the API server does not wait for confirmation from the kubelet.

Forced Pod termination is useful when Kubernetes pods are stuck in a terminating state.



Understanding the Pod Lifecycle



Duration: 10 mins

Problem Statement:

You've been tasked with understanding about the lifecycle of a Pod.

Assisted Practice: Guidelines

Steps to be followed:

1. Describing the Pod to understand its lifecycle



TECHNOLOGY

Working on Pod Allocation

Introduction

There are several ways to allocate Pods, and all the recommended approaches use label selectors to facilitate the selection.

The scheduler does a reasonable placement. However, users may sometimes want to control the node to which the Pod deploys. To do so, they can use any of the following methods:

- 01 nodeSelector field matching against node labels
- 02 Affinity and anti-affinity
- 03 nodeName field

nodeSelector

nodeSelector is the simplest recommended form of Node selection constraint.

Example of how to use nodeSelector:

Run **kubectl get nodes** command to get the names of the cluster nodes.

\$\$\$\$\$\$\$\$\$\$\$	\$\$\$:~ \$ ku	bectl get	nodes	
NAME	STATUS	ROLES	AGE	VERSION
k8s-head	Ready	master	4d	v1.11.2
k8s-node-1	Ready	<none></none>	4d	v1.11.2
k8s-node-2	Ready _	<none></none>	4d	v1.11.2



nodeSelector

Add a nodeSelector field to the Pod configuration.

For example:

```
apiVersion: v1
Kind: pod
Metadata:
name: nginx
Labels:
env: test
spec:
containers:
- name: nginx
image: nginx
```



nodeSelector

Add a **nodeSelector**.

```
apiVersion: v1
Kind: pod
Metadata:
name: nginx
Labels:
env: test
spec:
containers:
- name: nginx
image: nginx
imagePullPolicy: iFNotPresent
nodeSelector:
disktypr: ssd
```

Then, run **kubectl apply -f https://k8s.io/examples/pods/pod-nginx.yaml**. The Pod will get scheduled on the node that is attached to the label.

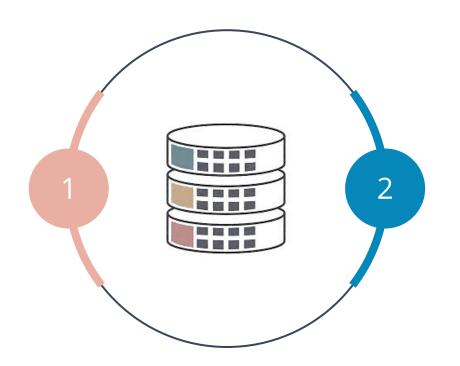


Node Isolation or Restriction

Node Isolation is used to ensure that only specific Pods run on nodes with certain isolation, security, or regulatory properties.

To make use of that label prefix for node isolation:

Use **Node authorizer** and enable **NodeRestriction** admission plugin



Add labels under the **node- restriction.kubernetes.io/**prefix to the nodes and use
those labels in the node
selectors

Affinity and Anti-affinity

The affinity and anti-affinity feature expands the types of constraints that users can define.

The key enhancements are:

The affinity or anti-affinity language is more expressive.

Users can constrain a Pod using labels on other Pods running on the node.

Users can indicate that the rule is **soft** or **preferred** rather than a hard requirement.



02

Node Affinity

Node Affinity is conceptually the same as nodeSelector. It allows constraining the nodes that the Pod is eligible to be scheduled on, based on the labels on the Node.

There are two types of node affinity:

01 requiredDuringSchedulingIgnoredDuringExecution

preferredDuringSchedulingIgnoredDuringExecution



Node Affinity

```
apiVersion: v1
 Kind: Pod
  Metadata:
   name: with-node-affinity
   spec:
    nodeAffinity:
     requiredDuringSchudulingIgnoredDuringExecution:
    nodeSelectorTerms:
     - matchExpressions:
     - key: kubernetes.io/e2e-az-name
       operator: in
        values:
         - e2e-az1
         - e2e-az2
           preferedDuringSchedulingDuringException:
      preference:
       matchExpressions:
     key: another-node-label-key
operator: In
       values:
         - another-node-label-value
        containers:
        - Name: with-node-affinity
image: k8s.gcr.io/pause:2.0
```

The Node Affinity rule says that the Pod can only be placed on a Node with a label whose key is **kubernetes.io/e2e-az-name** and the value is either **e2e-az1** or **e2e-az2**.

Node Affinity

The new node affinity syntax supports the following operators:

In, Notln, Exists, DoesNotExist, Gt, and Lt.

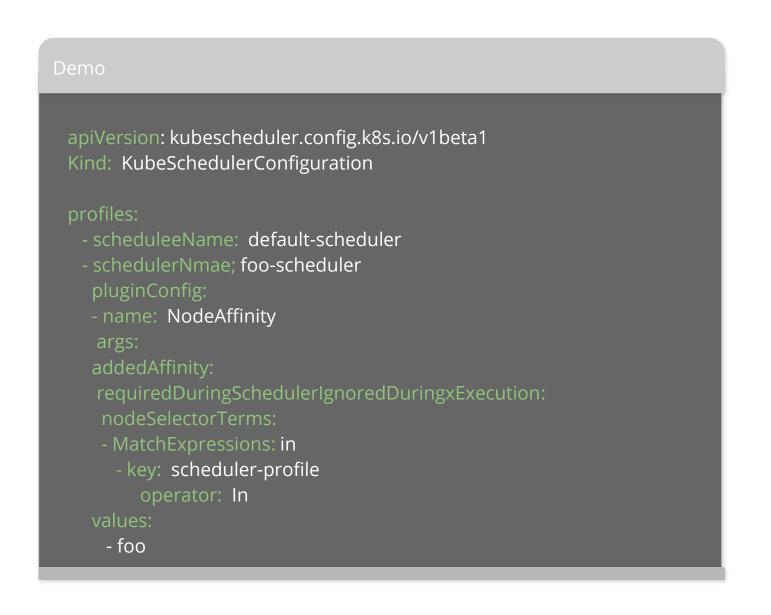
Both **nodeSelector** and **nodeAffinity** must be satisfied for the Pod to be scheduled on a candidate node.

If multiple **matchExpressions** associated with **nodeSelectorTerms** are specified, the Pod can be scheduled on a node only if all the **matchExpressions** are satisfied.



Feature State: Kubernetes V1.20 [Beta]

When configuring multiple scheduling profiles, associate a profile with a node affinity that is useful if a profile only applies to a specific set of n. Add an **addedAffinity** to the **args** field of the **NodeAffinity** plugin in the scheduler configuration. For example:





Feature State: Kubernetes V1.20 [Beta]

There are existing Kubernetes concepts that allow exposing a single service. This can be achieved through an Ingress by specifying a default backend with no rules:

```
apiVersion: networking.k8s.10/v1
Kind: Ingress
Metadata:
 name: test-ingress
 spec:
  defaultBackend
  service:
  name: test
  port: k8s.exaple.com
    number: 80
```



Inter-Pod Affinity and Anti-Affinity

Inter-pod affinity and anti-affinity allow constraining the nodes the Pod is eligible to be scheduled based on labels on Pods that are already running on the node, rather than on labels on the nodes.

There are two types of Pod affinity and anti-affinity:

01 requiredDuringSchedulingIgnoredDuringExecution

opening of the preferred During Scheduling Ignored During Execution



Inter-Pod Affinity and Anti-Affinity

```
apiVersion: v1
 Kind: pod
 Metadata:
   name: with-pod-affinity
   spec:
    affinity:
    nodeAffinity:
    requiredDuringSchudulingIgnoredDuringExecution:
     - labelSelector:
       matchExpressions:
       - key: security
        operator: in
        values:
         - 51
      topologyKey: topology.kebernetes.io/zone
    podAntiAffinity:
         preferedDuringSchedulingIgnoredException:
    - weight: 100
      podAffinityTerm:
      labelSelector
      matchExpressions:
    - key: security
     operator: In
       values:
        - 52
        topologyKey: topology.kubernetes.io/zone
     containers:
        - Name: with-pod-affinity
image: k8s.gcr.io/pause:2.0
```

An example of a Pod that uses Pod affinity:

The affinity on this Pod defines one Pod affinity rule and one Pod anti-affinity rule.



Inter-Pod Affinity and Anti-Affinity

The rules for Pod affinity and anti-affinity are as follows:

Pod affinity rule

The Pod can be scheduled on a node only if the node is in the same zone as at least one already-running Pod that has a label with key **security** and value **\$1**.

Pod anti-affinity rule

The Pod should not be scheduled on a node if the node is in the same zone as a Pod with label having key **security** and value **S2**.



Constraints on TopologyKey

For performance and security reasons, there are some constraints on **topologyKey**:



For Pod affinity, empty topologyKey is not allowed.



For Pod anti-affinity also, empty topologyKey is not allowed.



For **requiredDuringSchedulingIgnoredDuringExecution** Pod anti-affinity, the admission controller **LimitpodHardAntiAffinityTopology** was introduced to limit topologyKey to **kubernetes.io/hostname**.



Except for the above cases, the topologyKey can be any legal label key.

The affinity term is applied to the union of the namespaces selected by the namespaceSelector and the ones listed in the namespaces field.

An empty namespaceSelector ({}) matches all namespaces, while a null or empty namespaces list and null namespaceSelector matches the namespace of the Pod where the rule is defined.

Inter-pod affinity and anti-affinity can be more useful when they are used with higher level collections.

Users can easily configure that a set of workloads should be co-located in the same defined topology.

```
apiVersion: v1
Kind: Deployment
Metadata:
 name: redis-cache
 spec:
  selector:
  matchlabels:
  app: store
  replicas: 3
  template:
  metadata:
  labels:
    app: store
 spec:
   affinity:
    podAntiAffinity:
   requiredDuringSchudulingIgnoredDuringExecution:
   - labelSelector:
     matchExpressions:
     - key: app
      operator: in
      values:
       - store
    topologyKey: "kubernetes.io/hostname"
   containers:
       - Name: with-pod-affinity
       image: redis: 3.2- alpine
```

This example shows the YAML snippet of a simple Redis deployment with three replicas and selector label **app=store**.

The following example shows the YAML snippet of the webserver deployment that has Pod anti-affinity and Pod affinity configured:

```
apiVersion: apps/v1
 Kind: Deployment
  Metadata:
   name: web-server
   spec:
    selector:
    matchlabels:
    app: web-store
   replicas: 3
    template:
    metadata:
    labels:
      app: web-store
    spec:
    affinity:
    podAntiAffinity:
    requiredDuringSchudulingIgnoredDuringExecution:
     - labelSelector:
       matchExpressions:
       - key: app
```

```
operator: in
        values:
         - web-store
      topologyKey: "kubernetes.io/hostname"
    podAntiAffinity:
    requiredDuringSchudulingIgnoredDuringExecution:
     - labelSelector:
       matchExpressions:
       - key: app
        operator: in
        values:
         - store
        topologyKey: "kubernetes.io/hostname"
    containers:
        - Name: web-app
         image: nginx: 1.16- alpine
```

To create two Deployments, the three-node cluster should be:

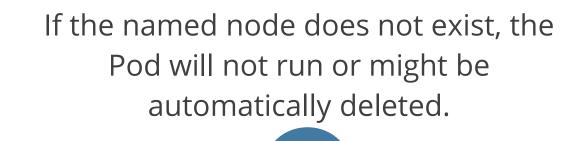
node-1	node-2	node-3
webserver-1	webserver-2	webserver-3
cache-1	cache-2	cache-3

All the three replicas of the web server are automatically co-located with the cache.

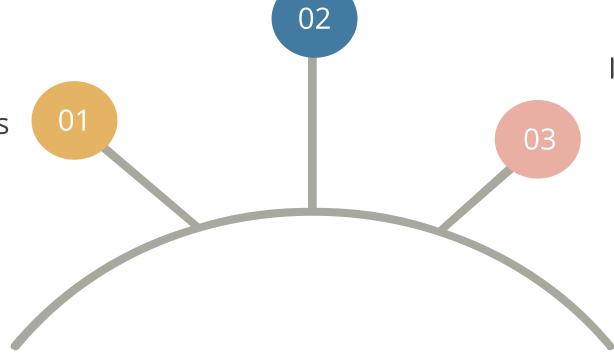


NodeName

nodeName is the simplest form of node selection constraint, but due to its limitations, it is rarely used. Some of the limitations of using **nodeName** are:



Node names in cloud environments are not always predictable or stable.



If the named node does not have enough resources, the Pod will fail with a reason indicated.

NodeName

An example of a Pod config file using the nodeName field is:

apiVersion: v1
Kind: pod
Metadata:
name: nginx
spec:
containers:
- name: nginx
image: nginx
nodename: kube-01

This Pod will run on the node **kube-01**.



Nodename and Nodeselector Affinity



Duration: 15 mins

Problem Statement:

You've been asked to assign Pods to nodes using nodeName and nodeSelector fields.

Assisted Practice: Guidelines

Steps to be followed:

- 1. Creating a Pod with the field nodeName
- 2. Creating a Pod with the field nodeSelector
- 3. Assigning Labels to nodes
- 4. Creating a Pod with the Notln operator



TECHNOLOGY

Init Containers

Introduction

Init containers are specialized containers that run before app containers in a Pod.



Init containers can contain utilities or setup scripts that are not present in an app image.



Init Containers

A Pod can have multiple containers running apps within it. It can also have one or more init containers that are run before the app containers are started.

The differences between regular and init containers are as follows:



Init containers support all the fields and features of app containers, including resource limits, volumes, and security settings.



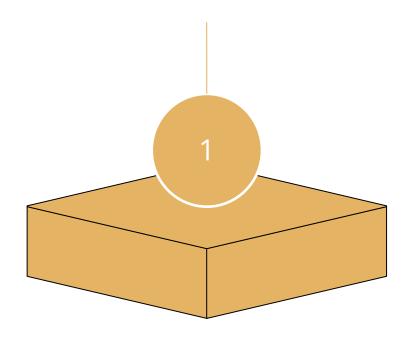
Init containers do not support **lifecycle**, **livenessProbe**, **readinessProbe**, **or startupProbe** because they must run to completion before the Pod can be ready.



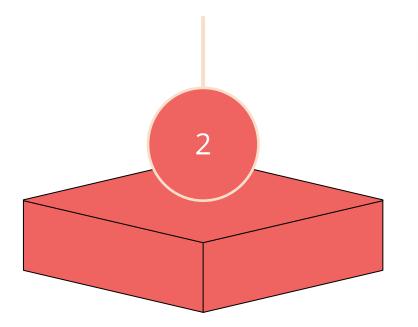
Init Containers

The advantages of init containers for start-up related code are:

Init containers can contain utilities or custom code for a setup that is not present in an app image.



The application image builder and deployer roles can work independently without the need to jointly build a single app image.

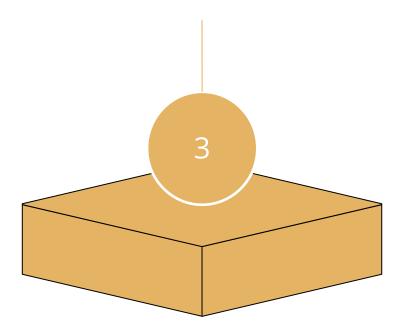


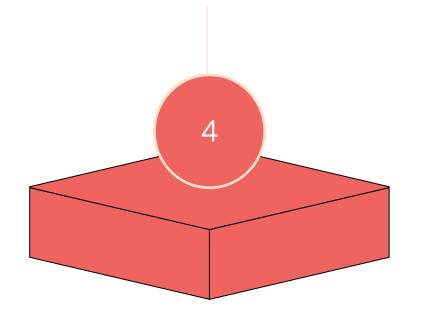
Init Containers

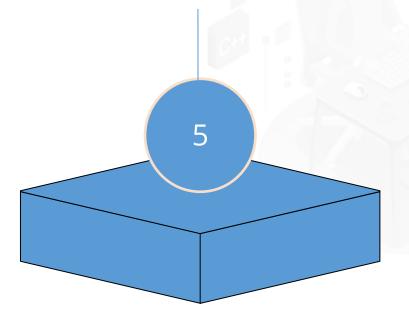
The advantages of init containers for start-up related code are:

Init containers offer a mechanism to block or delay app container startup until a set of preconditions are met.

Init containers can run with a different view of the filesystem than app containers in the same Pod. Init containers can securely run utilities or custom code that would otherwise make an app container image less secure.







This example defines a simple Pod that has two init containers:

```
apiVersion: v1
 Kind: pod
 Metadata:
   name: myapp-pod
   labels:
    app: myapp
   spec:
    containers:
    - name: myapp-container
     image: busybox:1.28
     command: ['sh','-c','echo The app os running! 66 sleep3600
    initContainers:
     - name: myapp-container
     image: busybox:1.28
     command: ['sh','-c',"intil nslookup
myservice.$(cat/var/run/secrets/kubernetes.io/seviceaccount/namespace).svc.cluster.l
ocal; do echo waiting for myservice; seep 2; done"]
    - name: myapp-container
     image: busybox:1.28
     command: ['sh','-c',"until nslookup
mydb.4(cat/var/run/secrets/kubernetes.io/serviceaccount/namespace.svc.cluster.local
;do echo waiting for mydb;seep2;done"]
```



Start the Pod by running:

kubectl apply -f myapp.yaml The output is similar to this: pod/myapp-pod created And check on its status with: kubectl get -f myapp.yaml The output is similar to this: NAME READY STATUS RESTARTS AGE myapp-pod 0/1 Init:0/2 0 6m or for more details: kubectl describe -f myapp.yaml



Here is a configuration that can be used to make the services appear:

```
apiVersion: v1
 Kind: service
 Metadata:
  name: myservice
  spec:
    ports:
   - protocol: TCP
     port: 80
     targetPort: 9376
apiVersion: v1
 Kind: service
 Metadata:
  name: mydb
  spec:
    ports:
   - protocol: TCP
     port: 80
     targetPort: 9377
```



To create the **mydb** and **myservice** services, use the given command:

kubectl apply -f services.yaml The output is similar to this: service/myservice created service/mydb created See the init containers complete, and that the myapp-pod pod moves into the Running state: kubectl get -f myapp.yaml The output is similar to this: READY STATUS RESTARTS AGE NAME myapp-pod 1/1 Running 0 9m



Using Init Containers



Duration: 15 mins

Problem Statement:

You've been asked to create a Pod using an init Container.

Assisted Practice: Guidelines

Steps to be followed:

- 1. Creating an Httpd Pod
- 2. Creating a Service
- 3. Verifying the Pods state



Using Multi-init Containers



Duration: 15 mins

Problem Statement:

You've been asked to create a Pod using multiple init Containers.

Assisted Practice: Guidelines

Steps to be followed:

- 1. Creating a Pod
- 2. Creating the first Service
- 3. Creating the second Service
- 4. Verifying the Pods state



TECHNOLOGY

Managing Container Resources

Introduction

The most common resources to specify in a Pod are CPU and memory (RAM).



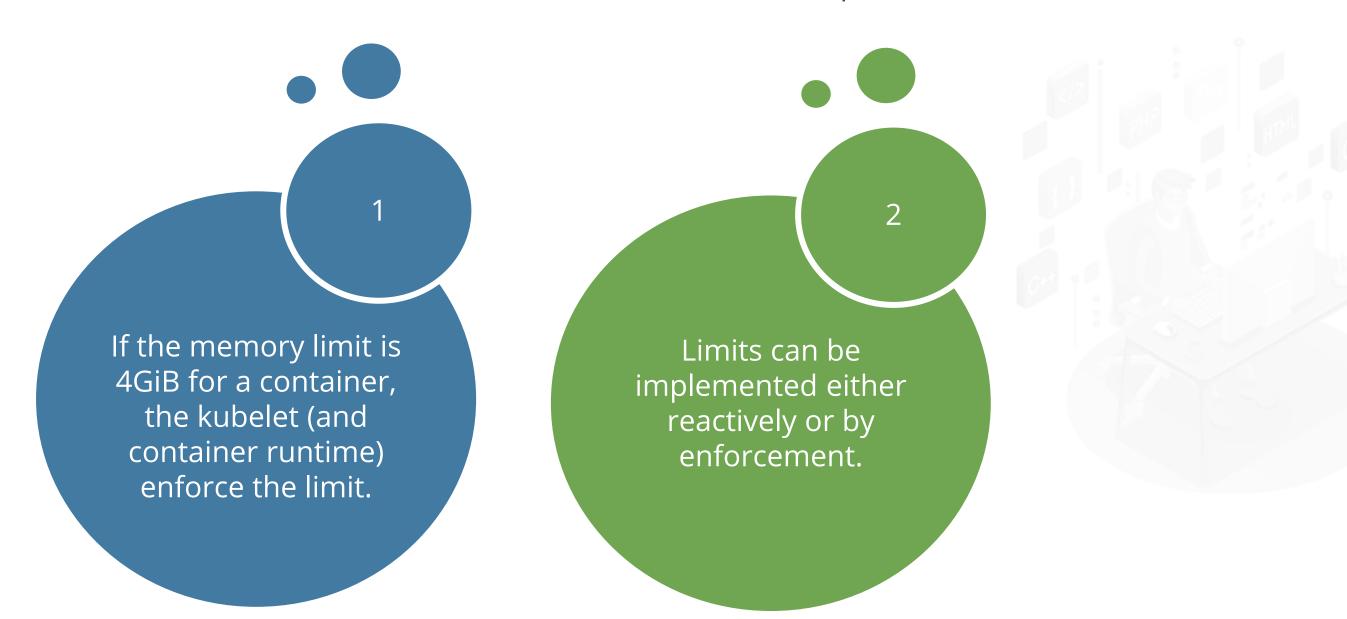
When the resource **request** is specified for containers in a Pod, the Kube-scheduler uses this information to decide on which node to place the Pod.



When the resource **limit** is specified for a container, the kubelet enforces those limits so that the running container is not allowed to use more of that resource than the limit set.

Requests and Limits

If the node where a Pod is running has enough resources available, it is possible (and allowed) for a container to use more resources than requested.



Resource Types

CPU and **memory** are both **resource types**.

CPU represents compute processing and is specified in units of Kubernetes CPUs.

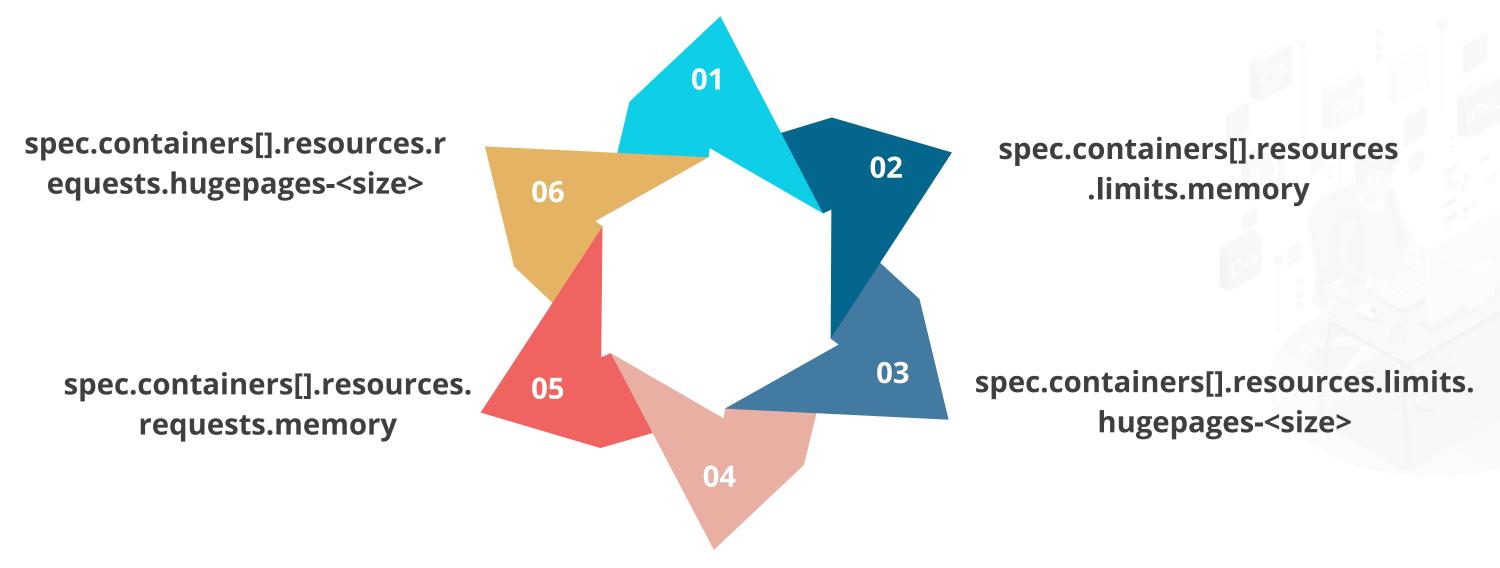


CPU and memory are collectively referred to as compute resources or resources.

Resource Requests and Limits of Pod and Container

Each container of a Pod can specify one or more of the following:





spec.containers[].resources.requests.cpu



Resource Units in Kubernetes

CPU resource units

In Kubernetes, one CPU is equivalent to one physical CPU core or one virtual core, depending on whether the node is a physical host or a virtual machine running inside a physical machine.

Memory resource units

Memory can be expressed as a plain integer or as a fixed-point number, using one of these suffixes:

E, P, T, G, M, K.

Resource Units in Kubernetes

```
apiVersion: v1
Kind: pod
Metadata:
 name: Frontend
 spec:
  containers:
  - name: app
    image: images.my company.example/app;V4
   resources:
      requests:
      memory: "64M1"
      cpu: '250M"
      limits:
      memory: "128M1"
      cpu: "500m"
   name: log-aggregator
    image: images.my company.example/log-aggregator;v6
    resources:
      requests:
      memory: "64M1"
      cpu: '250M"
      limits:
      memory: "128M1"
       cpu: "500m"
```

- In this example, the Pod has two containers.
- Each container has a request of 0.25 CPU and 64MiB (2²⁶ bytes) of memory.
- Each container has a limit of 0.5 to 1 CPU and 128MiB to 256 MiB of memory.

How Pods with Resource Limits are Run

When the kubelet starts a container of a Pod, it passes the CPU and memory limits to the container runtime.

If a container exceeds its memory limit, it might be terminated.

2

If a container exceeds its memory request, its Pod will likely be evicted whenever the node runs out of memory.

3

A container might or might not be allowed to exceed its CPU limit for an extended period.

Nodes have **local ephemeral storage**, backed by locally-attached writeable devices or, sometimes, by RAM.

Kubernetes supports two ways to configure local ephemeral storage on a node:

Single Filesystem

In this configuration, place the different kinds of ephemeral local data (emptyDir volumes, writeable layers, container images, logs) in one filesystem.

Two Filesystems

Use this filesystem for other data (for example, system logs that are not related to Kubernetes). It can even be the root filesystem.

Ephemeral storage can be used for managing local ephemeral storage. Each container of a Pod can specify one or more of the following:

spec.containers[].resources.limits.ephemeral-storage

spec.containers[].resources.requests.ephemeral-storage



In the following example, the Pod has two containers:

```
apiVersion: v1
Kind: pod
Metadata:
 name: Frontend
 spec:
  containers:
  - name: app
    image: images.my company.example/app;V4
   resources:
      requests:
       ephemeral-storage: "2Gi"
      limits:
       ephemeral-storage: "4Gi"
       cpu: "500m"
   - name: log-aggregator
    image: images.my company.example/log-aggregator;v6
    resources:
      requests:
       ephemeral-storage: "2Gi"
      limits:
       ephemeral-storage: "4Gi"
```



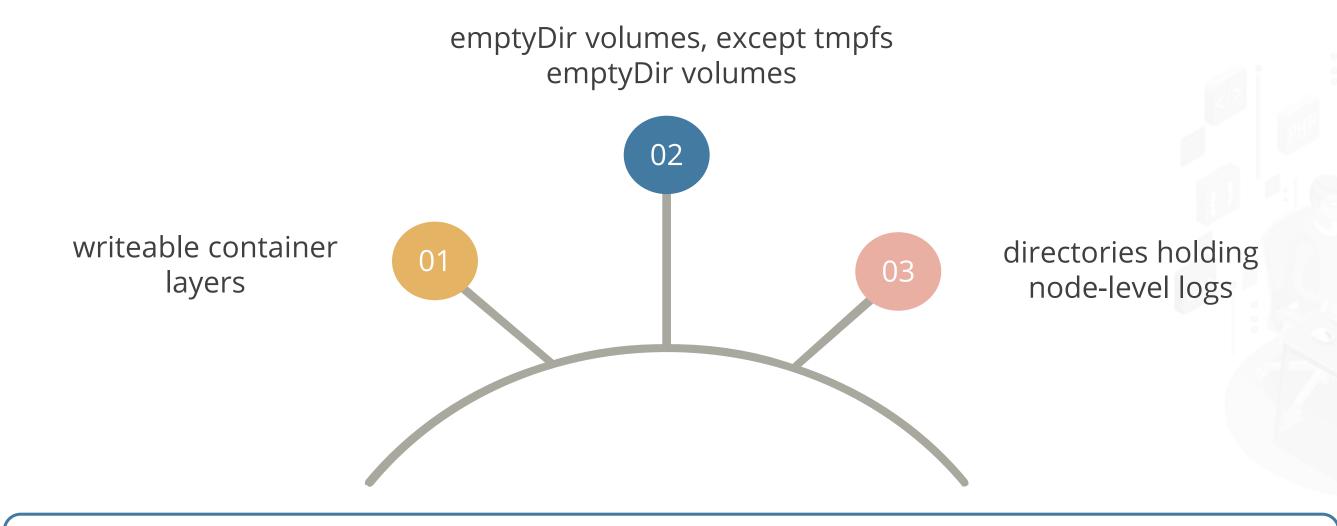
When a Pod is created, the Kubernetes scheduler selects a node on which the Pod will run.



The scheduler ensures that the sum of the resource requests of the scheduled containers is less than the capacity of the node.

Ephemeral Storage Consumption Management

If the kubelet is managing local ephemeral storage as a resource, it measures the use of storage in:

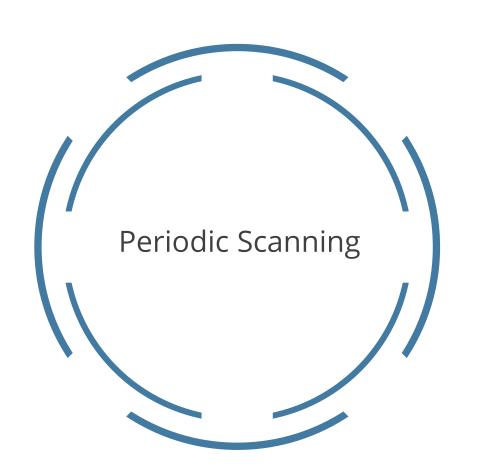


If a Pod uses more ephemeral storage than allowed, it receives an eviction signal from kubelet.



Periodic Scanning and File System Project Quota

The kubelet supports different ways to measure Pod storage use:





Container Resources



Duration: 10 mins

Problem Statement:

You've been asked to configure the memory requests and limits for the Containers.

Assisted Practice: Guidelines

Steps to be followed:

- 1. Creating a ResourceQuota
- 2. Defining the LimitRange
- 3. Verifying the ResourceQuota

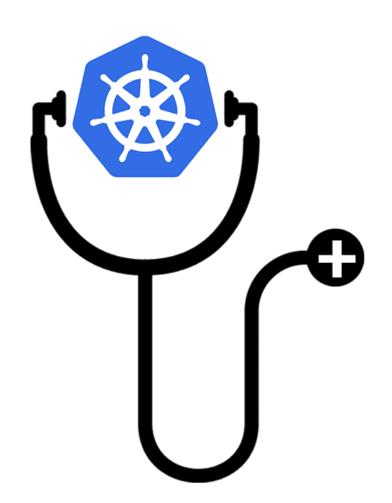


TECHNOLOGY

Health Monitoring

Overview

Node Problem Detector collects information about node problems from various daemons and reports these conditions to the API server as **NodeCondition** and **Event**.



Enabling Node Problem Detector

Kubectl provides the most flexible management of Node Problem Detector. For example:

```
apiVersion: apps/v1
 Kind: DaemonSet
 Metadata:
  name: node-problem-detector-v0.1
   namespace: kube-system
   labels:
    k8s-app: nude-problem-detector
    version: v0.1
    kubernetes.io/cluster-service: "true"
  spec:
   selector:
    matchlabels:
    k8s-app: nude-problem-detector
    version: v0.1
    kubernetes.io/cluster-service: "true
   template:
   metadata:
    labels:
    k8s-app: nude-problem-detector
    version: v0.1
    kubernetes.io/cluster-service: "true"
   spec:
    hostnetwork: true
    containers:
```

```
name: node-problem-detector
     image: k8s.gcr.io/node-problem-detector:v0.1
     securityContext:
      privileged: true
     resources:
      limits:
       cpu: " 200m"
       memory: "100M1"
       requests:
       cpu: "20m"
       memory: " 20M1"
     volumeMounts:
       - name: log
       mountpath: /log
       readOnly: true
     volumes:
      - name: log
       hostpath:
        path: /var/log/
```

The default configuration is embedded when building the Docker image of a Node Problem Detector. Use a **ConfigMap** to overwrite the configuration:

Change the configuration files in **config/**Create the ConfigMap node-problem-detector config:
kubectl create configmap node-problem-detector-config --from-file=config/



To change the **node-problem-detector.yaml**, use the **ConfigMap**:

```
apiVersion: apps/v1
 Kind: DaemonSet
 Metadata:
  name: node-problem-detector-v0.1
  namespace: kube-system
   labels:
    k8s-app: nude-problem-detector
    version: v0.1
    kubernetes.io/cluster-service: "true"
  spec:
   selector:
    matchlabels:
    k8s-app: nude-problem-detector
    version: v0.1
    kubernetes.io/cluster-service: "true
   template:
   metadata:
    labels:
    k8s-app: nude-problem-detector
    version: v0.1
    kubernetes.io/cluster-service: "true"
   spec:
    hostnetwork: true
    containers:
    - name: node-problem-detector
```

```
image: k8s.gcr.io/node-problem-detector:v0.1
privileged: true
     resources:
       limits:
        cpu: " 200m"
        memory: "100M1"
       requests:
        cpu: "20m"
        memory: " 20M1"
     volumeMounts:
       - name: log
        mountpath: /log
        readOnly: true
       - name: config # Overwrite the config/directory with ConfigMap
volume
        mountpath: /config
        readOnly: true
      volumes:
       - name: log
        hostpath:
         path: /var/log/
       - name: config # Define ConfigMap volume
        configMap:
         name: node-problem-detector-config
```

Recreate the Node Problem Detector with a new configuration file:

If you have a node-problem-detector running, delete before recreating kubectl delete –f hhtps://k8s.io/examples/debug/node-problem-detector,yanl kubectl apply –f hhtps://k8s.io/examples/debug/node-problem-detector-configmap,yanl

If a Node Problem Detector is running as a cluster Addon, overwriting a configuration is not supported.



Kernel Monitor

Kernel Monitor is a system log monitor daemon supported in the Node Problem Detector.

To support a new **NodeCondition**, create a condition definition within the **conditions** field in **config/kernel-monitor.json**, such as:

```
{
    "type": "NodeConditionType"
    "reason": "CameCaseDefaultNodeConditionReason"
    "message": "arbitrary Default node condition message"
}
```

To detect new problems, extend the **rules** field in **config/kernel-monitor.json** with a new rule definition:

```
{
  "type": "temporary/permanent"
  "condition": "NodeConditionOfPermanentIssue"
  "reason": "CameCaseShortReason"
  "message": "regexp matching the issue in the kernel log"
}
```

The kernel monitor uses the **translator** plugin to translate the internal data structure of the kernel log. Users can implement a new translator for a new log format.



Recommendations and Restrictions

When running the Node Problem Detector, expect an extra resource overhead on each node.

A resource limit is set for the Node Problem Detector.

The kernel log grows relatively slowly.

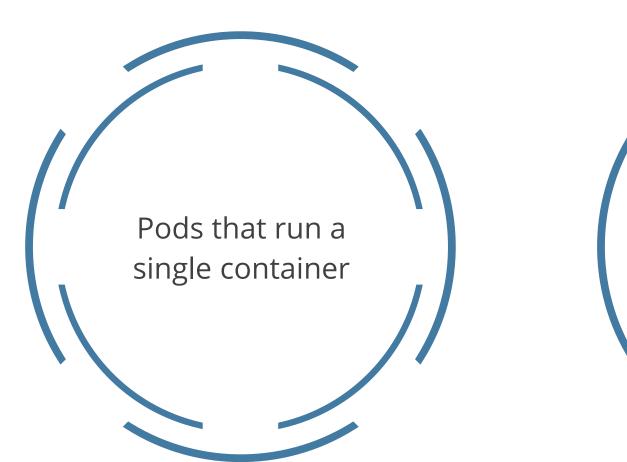
Even under high load, the resource usage is acceptable.

TECHNOLOGY

Multi-Container Pods

Using Pods in Different Situations

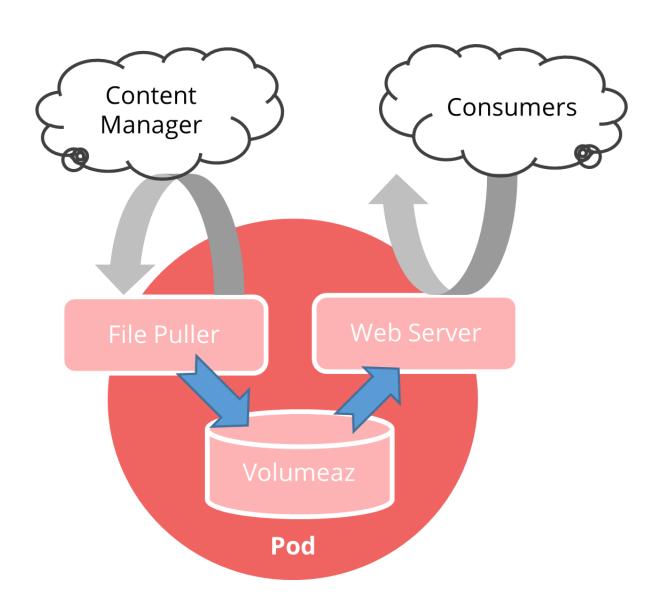
Pods are created using workload resources such as Deployment or Job. Pods in a Kubernetes cluster are used in two main ways:





Using Pods in Different Situations

Pods are designed to support multiple cooperating processes (as containers) that form a cohesive unit of service.



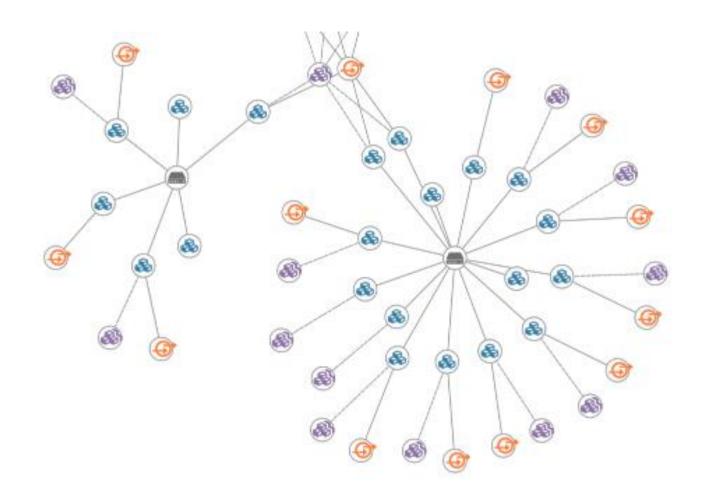
There can be a container that acts as a web server for files in a shared volume and a separate **sidecar** container that updates those files from a remote source.

TECHNOLOGY

Pod Topology Spread Constraints

Introduction

Topology spread constraints are used to control the Pods that are spread across the cluster, among failure domains such as regions, zones, nodes, and other user-defined topology domains.



It helps achieve high availability as well as efficient resource utilization.



Prerequisites

Topology spread constraints rely on node labels to identify the topology domain(s) that each Node is in. For example, a Node might have labels:

node=node1, zone=us-east-1a, region=us-east-1

Consider a four-node cluster with the following labels:

Name	Status	Roles	Age	Versions	Labels
Node1	Ready	<none></none>	4m26s	v1.16.0	node=node1,zone=zoneA
Node2	Ready	<none></none>	3m58s	v1.16.0	node=node2,zone=zoneA
Node3	Ready	<none></none>	3m17s	v1.16.0	node=node3,zone=zoneB
Node4	Ready	<none></none>	2m43s	v1.16.0	node=node4,zone=zoneB



Spread Constraints for Pods

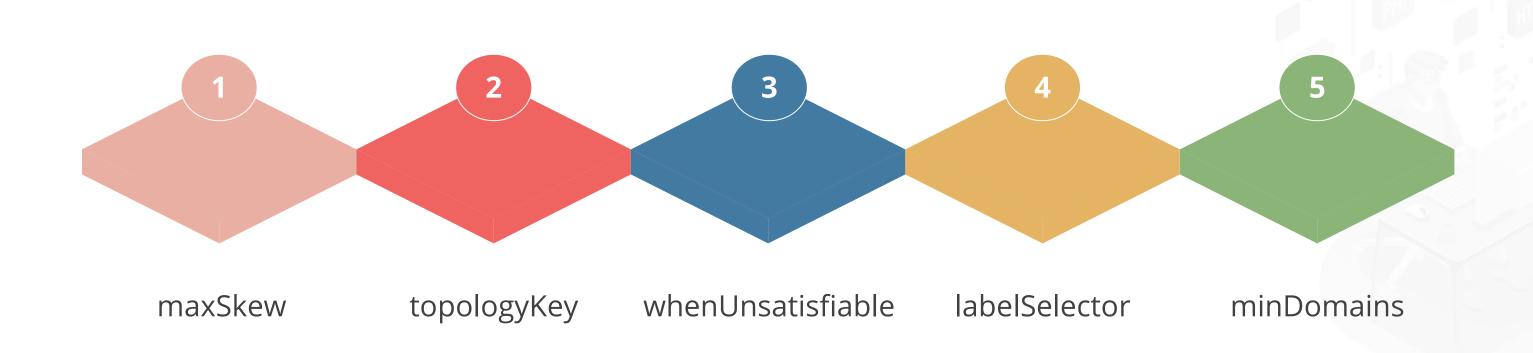
The API field **pod.spec.topologySpreadConstraints** is defined as below:

```
apiVersion: v1
Kind: pod
Metadata:
name: Mypod
spec:
topologySpreadConstraints:
maxSkew: <integer>
Topology: <string>
whenUnsatisfiable: <String>
labelSelector: <Object>
```



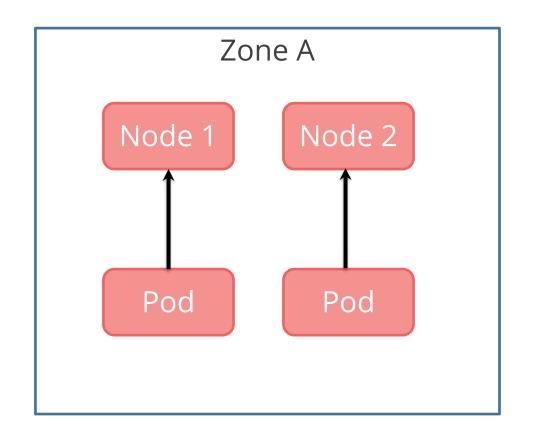
Spread Constraints for Pods

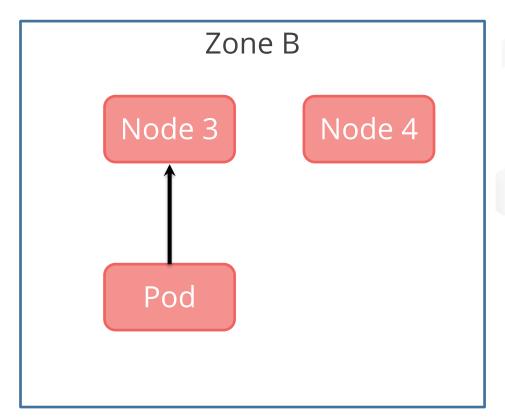
One or multiple **topologySpreadConstraint** can be defined to instruct the Kube-scheduler on how to place each incoming Pod to the existing Pods across the cluster. The fields are:



Example: One TopologySpreadConstraint

Consider a four-node cluster where three Pods labeled **foo:bar** are located in node1, node2, and node3 respectively:





Example: One TopologySpreadConstraint

For an incoming Pod to be evenly spread with existing Pods across zones, the spec can be given as:

```
Kind: Pod
apiVersion: v1
 Metadata:
  name: Mypod
        labels:
        foo: bar
 spec:
   topologySpreadConstraints:
    maxSkew: 1
        Topology: zone
        whenUnsatisfiable: DoNotSchedule
        labelSelector:
        matchLabels:
        foo:
        container:
        name: pause
        image: K8s.gcr.io/pause:3.1
```



Example: One TopologySpreadConstraint

The Pod spec can be created to meet various kinds of requirements:



Change **maxSkew** to a bigger value like **2** so that the incoming Pod can be placed on **zoneA** as well.



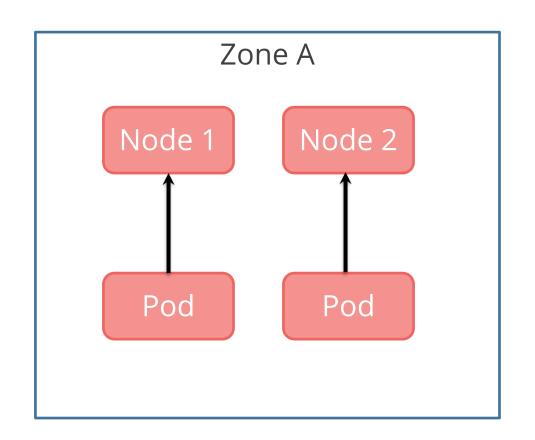
Change **topologyKey** to the **node** to distribute the Pods evenly across nodes instead of zones.

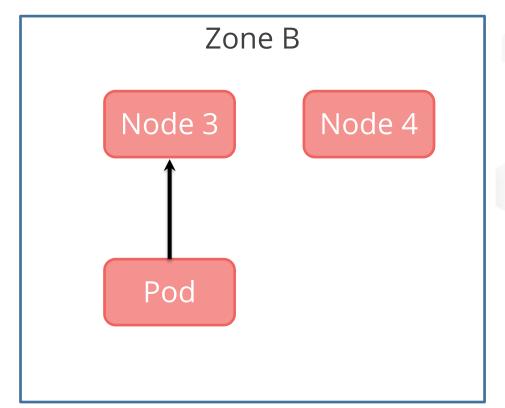


Change whenUnsatisfiable: DoNotSchedule to whenUnsatisfiable: ScheduleAnyway to ensure that the incoming Pod is always schedulable.

Multiple TopologySpreadConstraints

This builds upon the previous example. Consider a four-node cluster where three Pods labeled **foo:bar** are located in node1, node2, and node3 respectively:





Multiple TopologySpreadConstraints

Use two **TopologySpreadConstraints** to control the Pods spreading on both zone and node:

```
Kind: pod
apiVersion: v1
 Metadata:
   name: Mypod
        labels:
        foo: bar
 spec:
   topologySpreadConstraints:
    maxSkew: 1
        Topology: zone
        whenUnsatisfiable: DoNotSchedule
        labelSelector:
        matchLabels:
        foo: bar
         maxSkew: node
        whenUnsatisfiable: DoNotSchedule
        labelSelector:
        matchLabels:
        foo: bar
        container:
        name: pause
        image: K8s.gcr.io/pause:3.1
```



Conventions

There are some implicit conventions worth noting here:

Only the Pods holding the same namespace as the incoming Pod can be matching candidates.

Nodes without **topologySpreadConstraints[*].topologyKey** present will be bypassed.

Be aware of what will happen if the incomingPod's **topologySpreadConstraints[*].labelSelector** does not match its labels.

If the incoming Pod has **spec.nodeSelector** or **spec.affinity.nodeAffinity** defined, nodes not matching them will be bypassed.



Cluster-level Default Constraints

It is possible to set default Topology Spread Constraints for a cluster. Default Topology Spread Constraints are applied to a Pod only if:

It doesn't define any constraint in its .spec.topologySpreadConstraints.

It belongs to a service, replication controller, replica set, or stateful set.



Cluster-level Default Constraints

The following is an example of a configuration:

```
apiVersion: kubescheduler.config.k8s.io/v1betal
Kind: kubeschedulerconfiguration
Profile:
schedulerName: default-scheduler
Pluginconfig:
name: podTopologySpread
args:
defaultConstraints:
maxSkew: 1
TopologyKey: topology.kubernetes.io/zone
whenUnsatisfiable: ScheduleAnyway
defaultingType: List
```



Internal Default Constraints

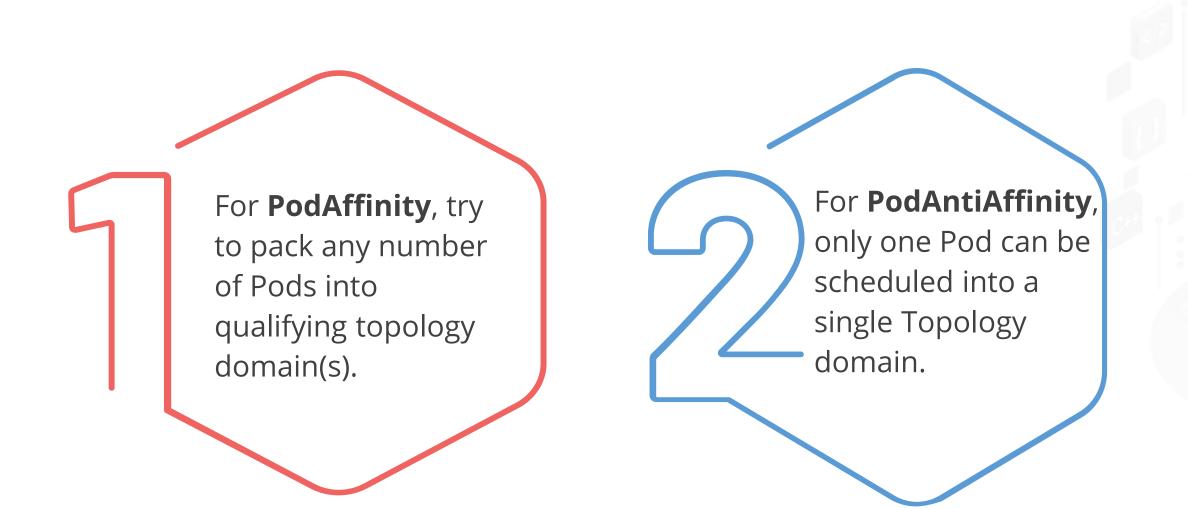
With the **DefaultpodTopologySpread** feature gate enabled by default, the legacy **SelectorSpread** plugin is disabled. Kube-scheduler uses the following default Topology Constraints for the **podTopologySpread** plugin configuration:

defaultConstraints:
 maxSkew: 3
 TopologyKey: "kubernetes.io/hostname"
 whenUnsatisfiable: ScheduleAnyway
 maxSkew: 5
 TopologyKey: "kubernetes.io/zone"
 whenUnsatisfiable: ScheduleAnyway



Comparison with PodAffinity and PodAntiAffinity

In Kubernetes, directives related to **Affinity** control how Pods are scheduled, more packed, or more scattered.



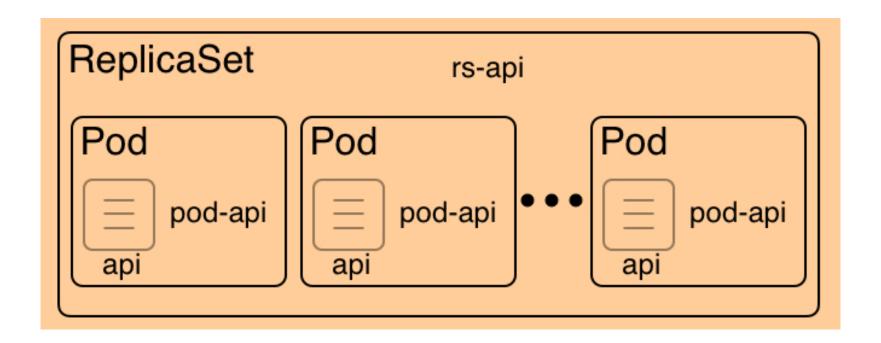
TECHNOLOGY

ReplicaSet

Purpose of ReplicaSet



A ReplicaSet's purpose is to maintain a stable set of replica Pods running at any given time. As such, it is often used to guarantee the availability of a specified number of identical Pods.



How a ReplicaSet Works

A ReplicaSet is defined with fields, including a selector that specifies how to identify Pods.

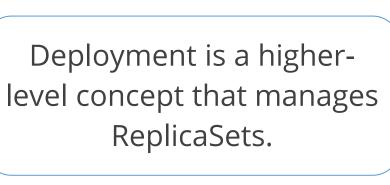
A ReplicaSet is linked to its Pods via the Pods' **metadata.ownerReferences** field, which specifies by what resource the current object is owned.

A ReplicaSet identifies new Pods to acquire by using its selector.

ReplicaSet in Use

A ReplicaSet ensures that a specified number of Pod replicas are running at any given time.

Deployment provides declarative updates to Pods along with a lot of other useful features.





Deployment is recommended over directly using ReplicaSets.



Using a ReplicaSet

Example:

```
apiVersion: app/v1
Kind: ReplicaSet
Metadata:
name: frontend
labels:
 app: guestbook
 tier: frontend
Spec:
#modify replicas according to your case
replicas: 3
selector:
matchLabers:
 tier: frontend
template:
metadata:
 labels:
 tier: frontend
Spec:
container:
name: php-redis
image: grc.io/google_samples/gb-frontened:v3
```



Using a ReplicaSet

To create the defined ReplicaSet and the Pods that it manages, save this manifest into the **frontend.yaml** and submit it to a Kubernetes cluster:

kubectl apply -f https://kubernetes.io/examples/controllers/frontend.yaml
get the current ReplicaSets deployed:
kubectl get rs

And see the frontend that is created:

NAME DESIRED CURRENT READY AGE
frontend 3 3 3 6s

Check on the state of the ReplicaSet:
kubectl describe rs/frontend



When to Use a ReplicaSet

The output will be similar to:

```
Name: frontend
Namespace: default
Selector: tier=frontend
Labels: app=guestbook
    tier=frontend
Annotations: kubectl.kubernetes.io/last-applied-
configuration:{"apiversion":"apps/v1","kind":"ReplicaSet","metadata":{"annotations":{},
"labels":{"app":"guestbook","tier":"frontend"},"name":"frontend",...
Replicas: 3 current/ 3 defired
pods Status: 3 running/0 waiting/ 0 succeeded / 0 failed
pod template:
 labels: tier=frontend
 containers:
 php-redis:
 image: gcr.io/google_sample/gb-frontend:v3
 port: <none>
 Host Port: <none>
 Environment: <none>
 Mounts: <none>
 volumes: <none>
Events:
```



When to Use a ReplicaSet

To check for the Pods brought up, use:

kubectl get pods

The Pod information will be shown like:

Name	Ready	Status	Restarts	Age
frontend-b2zdv	1/1	Running	0	6m36s
frontend-vcmts	1/1	Running	0	6m36s
frontend-wtsmm	1/1	Running	0	6m36s



When to Use a ReplicaSet

To verify that the owner reference of these Pods is set to the frontend ReplicaSet, get the yaml of one of the Pods running:

kubectl get Pods frontend-b2zdv -o yaml

apiVersion: app/v1

Kind: pod Metadata:

creationTimestamp: "2020-02-12T07:06:162"

generateName: frontend

labels:

tier: frontend

name: frontend-b2zdv Namespace: default ownerreferences: apiVersion: apps/v1

blockOwenersDeletion: true

controller: true kind: replicaSet name: frontend

uid: f391f6db-bb9b-4c09-ae74-6a1f77f3d5cf

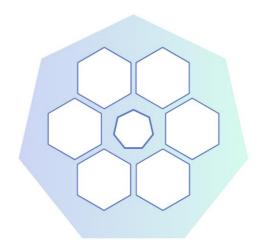
The output will be similar to the given image, with the frontend ReplicaSet's info set in the metadata's ownerReferences field.

To create bare Pods with no problems, it is strongly recommended that bare Pods do not have labels that match the selector of one of the ReplicaSets.



The reason for this is because a ReplicaSet is not limited to owning Pods specified by its template. It can acquire other Pods in the manner specified in the previous sections.

As the Pods do not have a Controller (or any object) as their owner reference and match the selector of the frontend ReplicaSet, they will be immediately acquired by it.



Suppose the Pods are created after the frontend ReplicaSet has been deployed and has set up its initial Pod replicas to fulfill its replica count requirement:

kubectl apply -f https://kubernetes.io/examples/pods/pod-rs.yaml



The new Pods will be acquired by the ReplicaSet and immediately terminated as the ReplicaSet would be over its desired count.

For fetching the Pods, use the given command:

kubectl get Pods

The output shows that the new Pods are either already terminated or are in the process of being terminated:

NAME	READY	STATUS	RESTARTS	AGE
frontend-b2zdv	1/1	Running	0	10m
frontend-vcmts	1/1	Running	0	10m
frontend-wtsmm	1/1	Running	0	10m
Pod1	0/1	Terminating	0	15
Pod2	0/1	Terminating	0	15



To create the Pods, run the following command:

kubectl apply -f https://kubernetes.io/examples/pods/pod-rs.yaml

After creating the Pods, create the ReplicaSet:

kubectl apply -f https://kubernetes.io/examples/controllers/frontend.yaml

Check that the ReplicaSet has acquired the Pods and has only created new ones according to its spec, until the number of its new Pods and the original Pods match its desired count. For fetching the Pods:

kubectl get pods



After running the **kubectl get pods** command, the output will be shown as:

Name	Ready	Status	Restarts	Age
frontend-hmmj2	1/1	Running	0	9s
pod1	1/1	Running	0	36S
pod2	1/1	Running	0	36S

Writing a ReplicaSet Manifest

As with all other Kubernetes API objects, a ReplicaSet needs the **apiVersion**, **kind**, and **metadata** fields. For ReplicaSets, the **kind** is always a ReplicaSet.



The name of a ReplicaSet object must be a valid **DNS subdomain name**.



A ReplicaSet also needs a **.spec** section.

Pod Template

PodTemplates are specifications for creating Pods and are included in workload resources such as Deployments, Jobs, and DaemonSets.



The **.spec.template** is a Pod template that is also required to have labels in place.



For the template's Restart Policy field, .spec.template.spec.restartPolicy, the only allowed value, is Always, which is the default.

Pod Selector

The label selector helps the users to identify a set of objects. It is the core grouping primitive in Kubernetes.

The **.spec.selector** field is a label selector. These are the labels used to identify potential Pods to acquire.

In the ReplicaSet, .spec.template.metadata.labels must match spec.selector, or it will be rejected by the API.



Replicas

Specify the number of Pods that should run concurrently by setting **.spec.replicas**. The ReplicaSet will create or delete its Pods to match this number.



If .spec.replicas is not specified, it sets to 1 by default.

Working with ReplicaSets

To delete a ReplicaSet and all its Pods, use **kubectl delete**. The Garbage Collector will automatically delete all the dependent Pods by default.



When using the REST API or the **client-go** library, ensure that **propagationPolicy** is set to **Background** or **Foreground** in the -d option.

Working with ReplicaSets

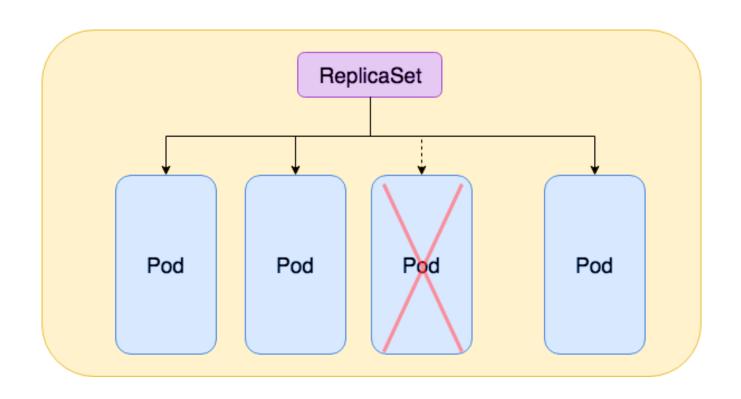
Example:

Kubectl proxy -port=800

Curl -x DELETE 'localhost:8080/apis/apps/v1/namespaces/default/replicasets/frontend'\

- >-d '{"kind":"DeleteOptions","apiVersion":"v1","propagationpolicy":"Foreground"}'\
- > -H "content-type: application/json

Working with ReplicaSets



Isolating Pods from a ReplicaSet

- Pods can be removed from a ReplicaSet by changing their labels.
- This technique may be used to remove Pods from Service for debugging, data recovery, and so on.
- Pods that are removed in this way will be replaced automatically.

Working with ReplicaSets

A ReplicaSet can be easily scaled up or down by simply updating the **.spec.replicas** field. When scaling down, the ReplicaSet controller prioritizes scaling down Pods by sorting the available Pods based on the following general algorithm:

1

Pending (and unschedulable) Pods are scaled down first.

2

If the **controller.kubernetes.io/Pod-deletion-cost** annotation is set, the Pod with the lower value will come first.

3

Pods on nodes with more Replicas come before Pods on nodes with fewer replicas.

4

If the Pods' creation times differ, the Pod that was created more recently comes before the older Pod.



ReplicaSet as a Horizontal Pod Autoscaler Target

A ReplicaSet can also be a target for **Horizontal Pod Autoscalers (HPA)**. Thus, a ReplicaSet can be auto-scaled by an HPA. Here is an example of HPA targeting the ReplicaSet:



ReplicaSet as a Horizontal Pod Autoscaler Target

Save this manifest into **hpa-rs.yaml** and submit it to a Kubernetes cluster to create the defined HPA that autoscales the target ReplicaSet, depending on the CPU usage of the replicated Pods:

kubectl apply -f https://k8s.io/examples/controllers/hpa-rs.yaml

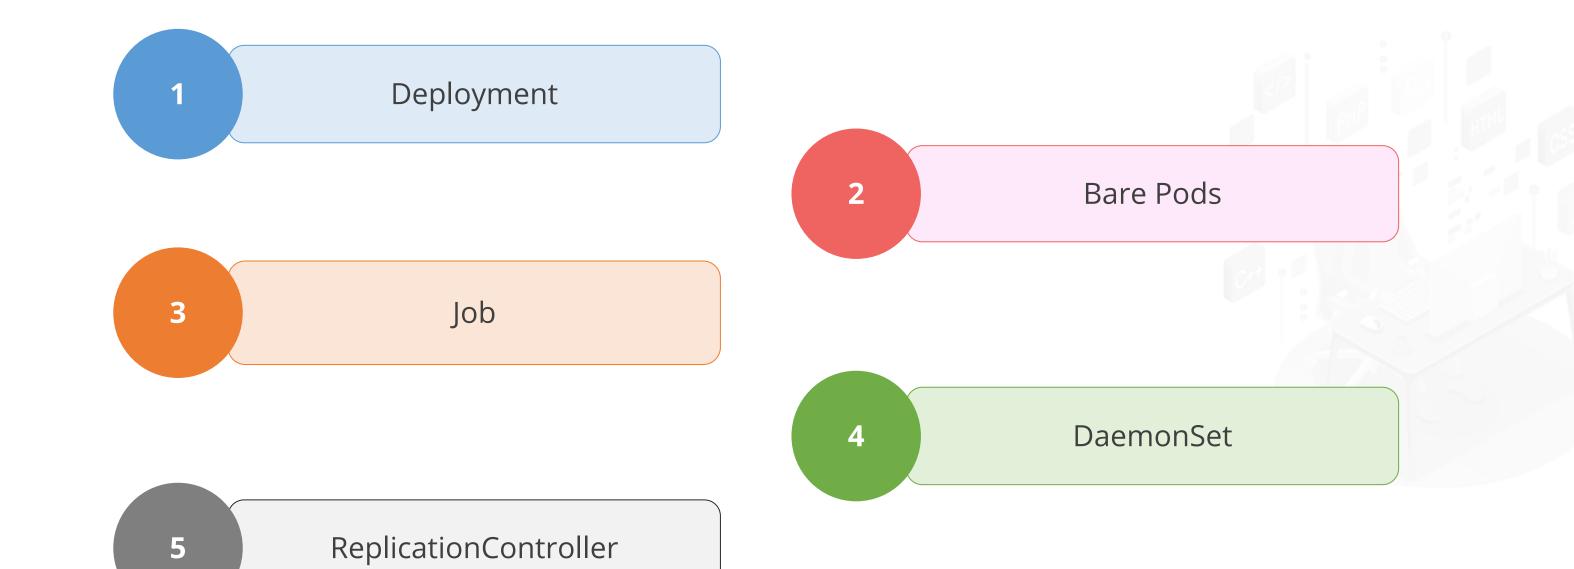
Alternatively, use the **kubectl autoscale** command to accomplish the same:

kubectl autoscale rs frontend --max=10 --min=3 --cpu-percent=50



Alternatives to ReplicaSet

The following is a list of alternatives to Kubernetes ReplicaSet:



ReplicaSets and Metrics Server



Duration: 20 mins

Problem Statement:

You've been asked to create a ReplicaSet and deploy a metrics server.

Assisted Practice: Guidelines

Steps to be followed:

- 1. Creating a ReplicaSet
- 2. Configuring the metrics server
- 3. Verifying the metrics server Deployment

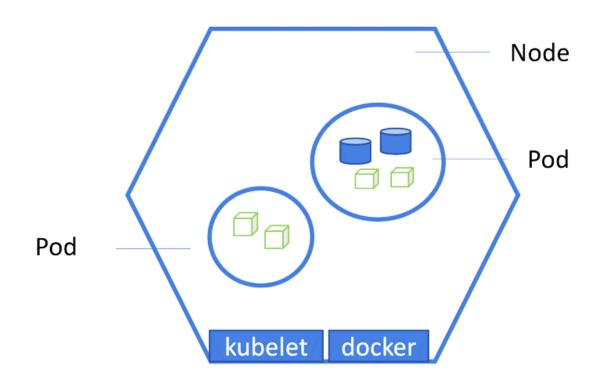


TECHNOLOGY

Static Pods

Overview

Static Pods are managed directly by the kubelet daemon on a specific node, without the API server observing them. They are always bound to one kubelet on a specific node.



- The main use of Static Pods is to run a self-hosted control plane.
- The kubelet automatically tries to create a mirror Pod on the Kubernetes API server for each static Pod.
- The Pods running on a node are visible on the API server but cannot be controlled from there.



Understanding Static Pods



Duration: 5 mins

Problem Statement:

You've been asked to create a static Pod.

Assisted Practice: Guidelines

Steps to be followed:

1. Create a static Pod in the worker node



TECHNOLOGY

Application Configuration

3

Overview

An application's configuration is everything that is likely to vary between deploys (staging, production, developer environments, and so on). This includes:

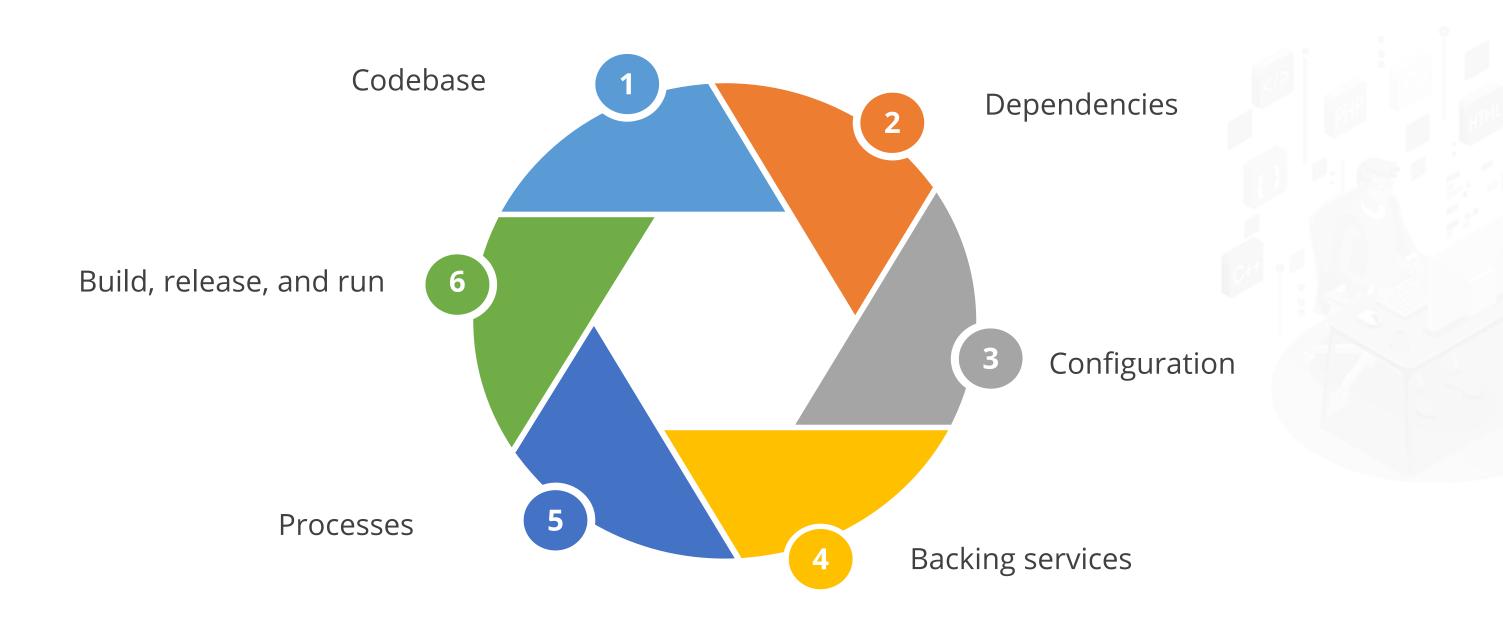
1 Resource handles to the database, Memcached, and other backing services

2 Credentials to external services such as Amazon S3 or Twitter

Per-deploy values such as canonical hostname for the deploy

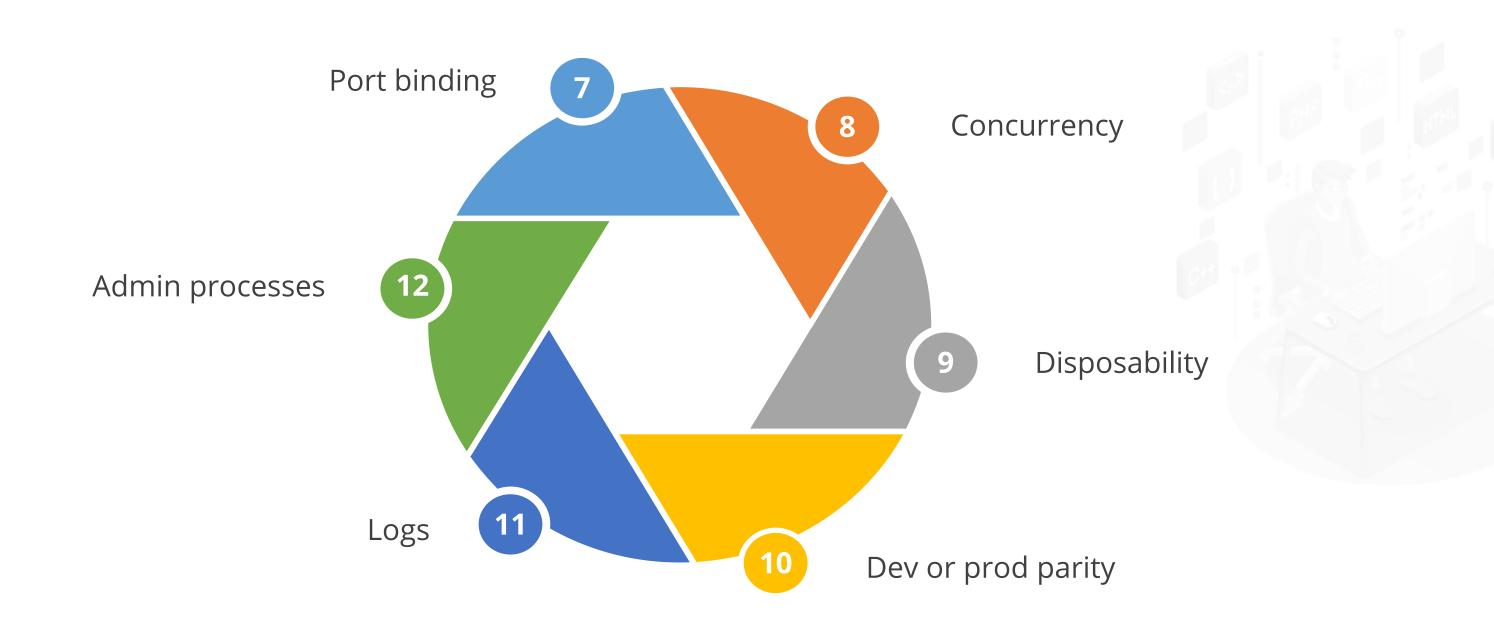
Twelve-Factor Principles

Kubernetes has evolved and is heavily influenced by twelve-factor principles:



Twelve-Factor Principles

Kubernetes has evolved and is heavily influenced by twelve-factor principles:



A Typical Way for Configuration in Kubernetes

To configure the apps in Kubernetes, use:



Using Environment Variables for Configuration

Here is an example of how to use environment variables by specifying them directly within the Pod specification:

```
apiVersion: v1
Kind: pod
Metadata:
        name: pod1
Spec:
        containers:
                 name: nginx
                 image: nginx
                          name: ENVVAR1
                          value: value 1
                          name: ENVVAR2
                           value: value2
```



Using Environment Variables for Configuration

Define two variables **ENVVAR1** and **ENVVAR2** in **spec.containers.env** with the values **value1** and **value2** respectively.

Use the **kubectl apply** command to submit the Pod Information to Kubernetes:

Use grep to filter the variable(s):

\$ kubectl apply -f
https://raw.githubusercontent.com/ab
hirockzz/kubernetes-in-anutshell/master/configuration/kinconfig-envvar-in-pod.yaml
pod/pod1 created

\$ kubectl exec pod1 -it -- env | grep ENVVAR ENVVAR1=value1 ENVVAR2=value2

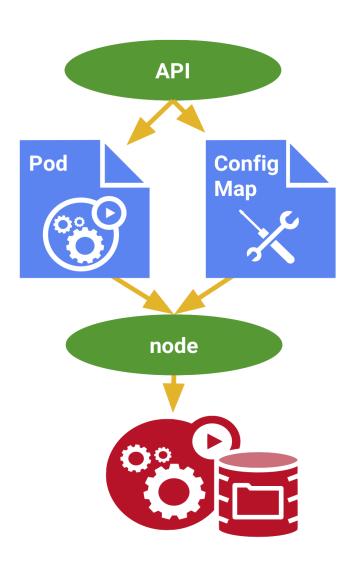
TECHNOLOGY

©Simplilearn. All rights reserved.

ConfigMap and Secrets

ConfigMaps

A ConfigMap is an API object used to store non-confidential data in key-value pairs. Pods can consume ConfigMaps as environment variables, command-line arguments, or as configuration files in a volume.



- It allows decoupling of environmentspecific configuration from the container images so that the applications are easily portable.
- It is not designed to hold large chunks of data. The data stored in a ConfigMap cannot exceed 1 MiB.

ConfigMap Object

A ConfigMap is an API object that lets users store configurations for other objects to use. Unlike most Kubernetes objects that have a **spec**, a ConfigMap has **data** and **binaryData** fields.

The **data** field is designed to contain UTF-8 strings, while the **binaryData** field is designed to contain binary data as base64-encoded strings.

Each key under the data or the binaryData field must consist of alphanumeric characters, hyphen (-), underscore (_), or period (.).

The keys that are stored in the **data** must not overlap with the keys in the **binaryData** field.



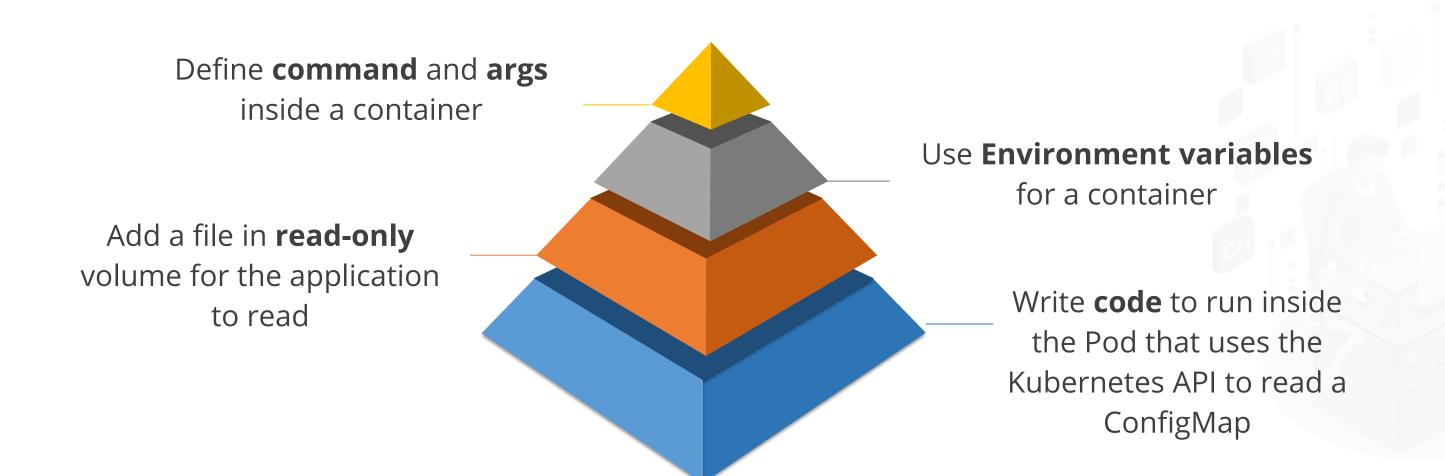
ConfigMaps and Pods

apiVersion: v1 Kind: configmap Metadata: name: game-demo data: # property-like keys; each key map to a simple value player_initial_lives: 3 ui_properties_file_name: "user-interface.properties" #file-like keys game.properties: enemy.types=aliens, monster players.maximum-lives=5 user-interface.properties:: color.good=purple1 color.bad=yellow allow.textmode=true

- The users can create a Pod specification that refers to a ConfigMap.
- The spec can configure the container(s) in that Pod based on the ConfigMap data.
- The Pod and the ConfigMap must be in the same namespace.
- The example shows a ConfigMap that has some keys with single values and other keys where the value looks like a fragment of a configuration format.

ConfigMaps and Pods

There are four different ways that a **ConfigMap** is used to configure a container inside a Pod:





Using ConfigMaps

ConfigMaps can be mounted as data volumes. They may also be used by other parts of the system, without having a direct exposure to the Pod.



Using ConfigMaps

To use ConfigMaps as files from a Pod to consume a ConfigMap in a volume in a Pod:



Create a ConfigMap or use an existing one



Modify the Pod definition to add a volume under .spec.volumes[]



Add a .spec.containers[].volumeMounts[] to each container that needs the ConfigMap



Modify the image or command line so that the program looks for files in that directory



Using ConfigMaps

Each ConfigMap used needs to be referred to in .spec.volumes.

If there are multiple containers in the Pod, each container needs its own **volumeMounts** block, but only one **.spec.volumes** is needed per ConfigMap.

Mounted ConfigMaps are updated automatically. When a ConfigMap currently consumed in a volume is updated, projected keys are updated as well.

The kubelet checks whether the mounted ConfigMap is fresh on every periodic sync. However, the kubelet uses its local cache for getting the current value of the ConfigMap.

Immutable ConfigMaps

This feature is controlled by the **ImmutableEphemeralVolumes** feature gate. An immutable ConfigMap can be created by setting the **immutable** field to **true**.



Once a ConfigMap is marked as immutable, it is not possible to revert the change or mutate the contents of the **data** or the **binaryData** field.

Secrets

Kubernetes Secrets help store and manage sensitive information such as passwords, OAuth tokens, and SSH keys. To use a Secret, a Pod needs to reference the Secret. A Secret can be used with a Pod in three ways:

As files in a volume mounted on one or more of its containers

As container environment variables

When images are pulled for the Pod by the kubelet



Types of Secrets

Kubernetes does not impose any constraints on the type name. However, if a built-in type is being used, all the requirements defined for that type must be met. The types are as follows:

Opaque secrets

SSH authentication secrets

Service account token secrets

TLS secrets

Docker config secrets

Bootstrap token secrets

Basic authentication secret



Types of Secrets

When creating a Secret, specify its type using the type field of the Secret resource, or certain equivalent kubectl command-line flags. These types vary in terms of the validations performed and the constraints Kubernetes imposes on them:

Built-in Type	Usage
Opaque	arbitrary user-defined data
kubernetes.io/service-account-token	service account token
kubernetes.io/dockercfg	serialized ~/.dockercfg file
kubernetes.io/dockerconfigjson	serialized ~/.docker/config.json file
kubernetes.io/basic-auth	credentials for basic authentication
kubernetes.io/ssh-auth	credentials for SSH authentication
kubernetes.io/tls	data for a TLS client or server
bootstrap.kubernetes.io/token	bootstrap token data



Creating a Secret

There are several options to create a Secret:

Use kubectl command
 Use config file
 Use kustomize

Editing a Secret

Example

please edit the object below. Lines beginning with a # will be ignored, and an empty file will abort the edit. If an error occurs while saving this file will #reopened with the relevant failure apiVersion: v1 data: username: YWRtaW4= password: MWYyZDFlMmU2NRm Kind: Secret Metadata: annotations: game-demo kubectl.kubernetes.io/last-applied-configuration: {..} creationTimestamp: 2016-01-22T18:41:56Z name: mysecret namespace: default resourceVersion: "164619" uid: cfee02d6-c137-11e5-8d73-42010af00002 Type: Opaque

An existing Secret may be edited with the command:

kubectl edit secrets mysecret

This opens the default editor and allows to update the base64 encoded Secret values in the **data** field.



How to Use Secrets

Mount Secrets as data volumes or expose them as environment variables that can be used by a container in a Pod. Other parts of the system can also use Secrets, without being directly exposed to the Pod. To access data from a Secret in a Pod:



Use a pre-existing Secret or create a new one



Change the Pod definition to add a volume under .spec.volumes[]



Specify a **.spec.containers[].volumeMounts[]** to each container that needs the Secret



Alter the command line or image so that the program looks for files in the specific directory



Usage of Secrets as Environment Variables

To use a Secret in a Pod's environment variable:

Create a Secret or use an existing one. Multiple Pods can reference the same Secret

Modify the image or command line so that the program looks for values in the specified environment variables



Modify the Pod definition in each container. The environment variable that consumes the Secret key should populate the Secret's name and key in env[].valueFrom.secretKeyRef

Using Secrets

An example of a Pod that uses Secrets from environment variables:

apiVersion: v1 Kind: pod Metadata: name: secret-env-pod Spec: containers: name: mycontainer image: redis name: SECRET-USERNAME valueFrom: secretKeyRef: name: mysecret key: username name: SECRET_PASSWORD valueFrom: secretKeyRef: name: mysecret key: password Restartpolicy: Never



Using Secrets

The Secret keys appear as normal environment variables inside a container that consumes a Secret in the environment variables. The following is the result of commands executed inside the container:

echo \$SECRET_USERNAME

echo \$SECRET_PASSWORD

The output is like:

Admin

1f2d1e2e67df

Environment variables are not updated after a Secret update.



Immutable Secrets and ConfigMaps

The Kubernetes feature **Immutable Secrets and ConfigMaps** provides an option to set individual Secrets and ConfigMaps as immutable:



Understanding Config Maps and Secrets



Duration: 10 mins

Problem Statement:

You've been asked to create secrets using kubectl by adding a config map.

Assisted Practice: Guidelines

Steps to be followed:

- 1. Add a config map entry to the Pod
- 2. Create secrets using kubectl



TECHNOLOGY

Achieving Scalability

Scaling Applications

Scalability is one of the core benefits of Kubernetes. The workloads can be automatically scaled up or down, depending on historical resource utilization with autoscaling. Autoscaling in Kubernetes has three dimensions:

Horizontal Pod Autoscaler (HPA)

Adjusts the number of replicas of an application

Cluster Autoscaler

Adjusts the number of nodes of a cluster

Vertical Pod Autoscaler (VPA)

Adjusts the resource requests and limits of a container



Scaling Applications

The autoscalers operate on one of two Kubernetes layers:

Pod level

The **HPA** and **VPA** methods take place at the Pod level. Both HPA and VPA will scale the available resources or instances of the container.

Cluster level

The **cluster autoscaler** falls under the cluster level, where it scales up or down the number of nodes inside the cluster.

Overview

When deploying a horizontally scalable application to Kubernetes, ensure the following configurations:



Resource Requests and Limits

Resource requests and limits dictate how much memory and CPU the application may use.

estimate the capacity and ensure that the cluster expands (and contracts) as the demand changes.

If there are no limits, the Kubernetes **scheduler** cannot ensure that workloads are spread evenly across the nodes.



It is vital to set limits and requests on a **Deployment**.

Node and Pod, Affinities and Anti-affinities

Affinities and anti-affinities enable users to distribute or restrict workload based on a set of rules or constraints.

The **scheduler** can use affinities and anti-affinities to assign the best node.

2

There are two kinds of affinity or anti-affinity: **preferred** and **required**.

The Pods can be placed according to specific requirements, with a combination of two types of affinity.



Health Checks

Kubernetes has two types of health checks:

Readiness Probe Notifies Kubernetes that the Pod is ready to start receiving requests

Liveness Probe Notifies Kubernetes that the pod is still running as expected



Configuring Liveness Probes



Duration: 5 mins

Problem Statement:

You've been asked to create and configure Pod using liveness probes.

Assisted Practice: Guidelines

Steps to be followed:

- 1. Create a Pod using liveness probes
- 2. Describe the Pod



Deployment Strategies

Deployment strategies dictate how the replacement of running Pods is done on Kubernetes when configuration updates are happening. There are two kinds of strategies:

Recreate strategy kills all the Pods managed by the Deployment, prior to starting a new Pod.

Rolling update strategy runs old as well as new configurations alongside each other.

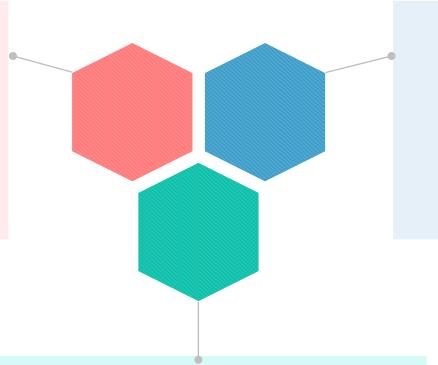


Pod Disruption Budgets

Disruptions in Kubernetes clusters are the norm, not the exception.

Reasons for disruptions:

A spot instance on AWS must be taken out of service.



An autoscaler removes any underutilized node.

The infrastructure team must apply a new config so that they can start replacing nodes within a cluster.



Pod Disruption Budgets

The job of the Pod disruption budget is to ensure that there is always a minimum number of Pods ready for Deployment.

With Pod disruption budget, it is possible to specify **minAvailable** or **maxUnavailable**.

2

The values may be a percentage of the desired replica count or an absolute number.

With Pod disruption budget, Kubernetes can reject **voluntary disruptions**.



Horizontal Pod Autoscalers

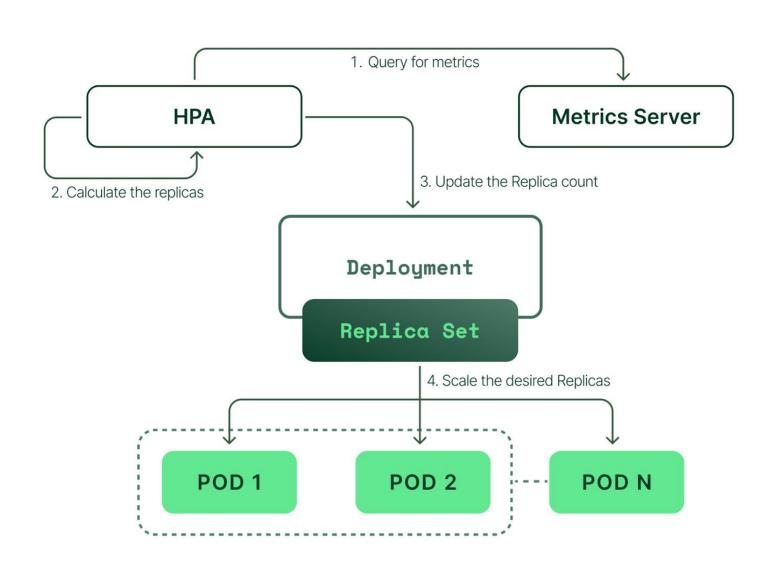
Kubernetes control plane can scale applications based on their resource utilization. As Pods become busier, the control plane brings up new replicas to share the load automatically.

Periodically, the controller monitors metrics provided by the metrics-server to scale the deployment up or down based on the Pods' load.

It is possible to scale while configuring a Horizontal Pod Autoscaler based on CPU and memory usage. But custom metrics servers can help extend the metrics available.

Using HPA for Dynamic Workload Management

HPA works in a **check, update, and check again** style loop.



Each of the steps in the loop works as follows:

- 1. Monitors the metrics server for resource usage
- 2. Calculates the desired number of replicas required based on the collected resource usage
- 3. Decides to scale up the application to the desired number of replicas
- 4. Changes the desired number of replicas

Since HPA is continuously monitoring, the process repeats from Step 1.



Understanding Horizontal Pod Auto Scaling



Duration: 10 mins

Problem Statement:

You've been asked to create a horizontal Pod autoscaler.

Assisted Practice: Guidelines

Steps to be followed:

- 1. Create an HPA in the master node
- 2. Check the Deployment Service
- 3. Verify HPA



TECHNOLOGY

Building Self-Healing Pods with Restart Policies

Introduction: No Time for Downtime

Any application that does not operate 24-7 is considered inefficient. The applications are expected to run uninterrupted, even if there are technical glitches, feature updates, or natural disasters.



Kubernetes facilitates the smooth working of the application by abstracting machines physically. It is a container orchestration tool. The Pods and containers in Kubernetes are self-healing.

The Three Container States-1

Waiting: Created but not running

A container in the waiting stage can still run operations such as pulling images or applying Secrets. To check the status of waiting, use:

kubectl describe pod [podname]

Along with the Pod's state, a message and reason for the state are displayed as:

• • •

State: Waiting

Reason: ErrImagePull

•••

The Three Container States-2

Running:

These are the containers running without issues. The given command is run before the Pod enters the running state:

postStart

Running pods display the time of starting of the container:

•••

State: Running

Started: Sat, 30 Jan 2021 16:48:38 +0530

•••

The Three Container States-3

Terminated:

It is a container that has failed or completed its execution. The given command is executed before the Pod is moved to be terminated:

preStop

Terminated Pods will display the time of the entrance of the container:

•••

State: Terminated Reason: Completed

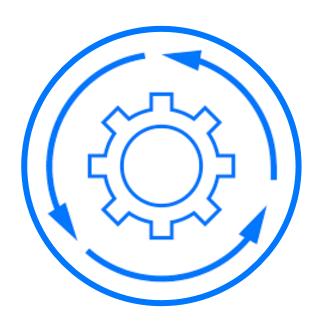
Exit Code: 0

Started: Mon, 11 Jan 2021 12:23:45 +0530 Finished: Mon, 18 Jan 2021 21:32:54 +0530

•••

Self-Healing in Kubernetes

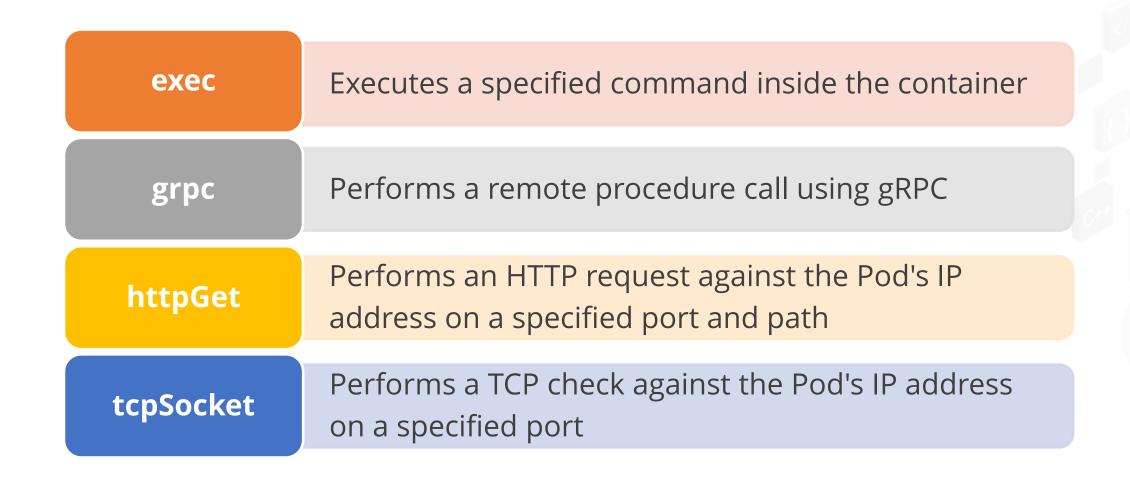
The Pod phase in Kubernetes offers insights into the Pod's placement. It can have:



- Pending Pods: The container is waiting to be started.
- Running Pods: The container has been created, and processes are running in it.
- **Succeeded Pods:** The Pods have completed the container life cycle.
- **Failed Pods:** The container either exited with non-zero status or was terminated by the system.
- Unknown Pods: The state of the Pod could not be obtained.

Kubernetes' Self-Healing

There are four distinct ways to inspect a container using a probe. Each probe must specify one of the following four mechanisms:



Kubernetes' Self-Healing Concepts

Each probe provides one of three results:

Failure

The container failed the diagnostic.

Unknown

The diagnostic failed, so no action can be taken.

Success

The container passed the diagnostic.



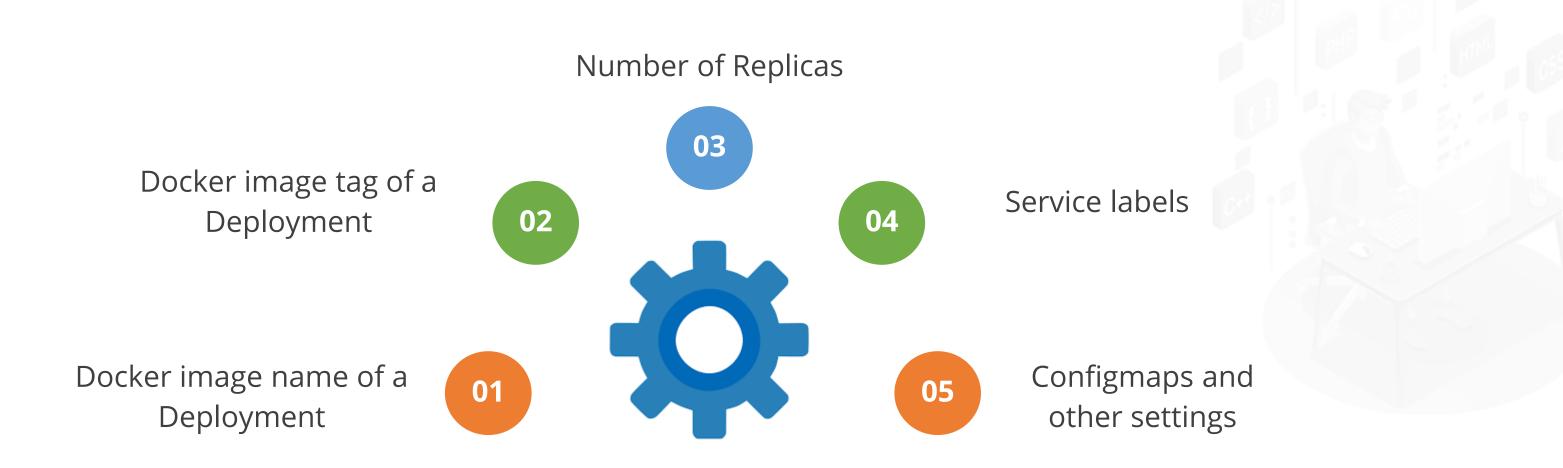
TECHNOLOGY

Manifest Management and Common Templating Tools

Overview

There is a need to use templates in Kubernetes manifests for common parameters.

These include:



Overview

Kubernetes does not have any default templating mechanism. Deployed manifests are static YAML files. If the parameters need to pass in the manifests, an external solution is needed.

Helm is the package manager for Kubernetes. It has templating capabilities. It is the Kubernetes equivalent of **yum** or **apt**.

Helm deploys charts. It is a collection of all pre-configured, versioned application resources that may be deployed as a single unit.

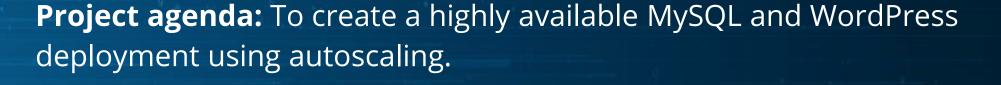
Key Takeaways

- Node Isolation is used to ensure that only specific Pods run on nodes with certain isolation, security, or regulatory properties.
- Topology spread constraints are used to control the Pods that are spread across the cluster, among failure domains such as regions, zones, nodes, and other user-defined topology domains
- A ReplicaSet ensures that a specified number of Pod replicas are running at any given time.
- Kubernetes control plane can scale applications based on their resource utilization. As Pods become busier, it automatically brings up new replicas to share the load.



MySQL and WordPress Installation in Kubernetes

Duration: 30 Min



Description:

Your team is going to deploy MySQL and WordPress containers. All the users should be added using ConfigMaps and all the sensitive data should be added using Secrets. The service should be on the NodePort. The WordPress Pod should not deploy if the MySQL service is not deployed. Also, ensure autoscaling is on.



MySQL and WordPress Installation in Kubernetes





Steps to perform:

- . Creating a cluster
- 2. Creating a Secret for storing the MySQL password securely
- 3. Creating a MySQL YAML manifest file for deploying MySQL
- 4. Deploying the WordPress application and designing the manifest YAML file
- 5. Accessing the WordPress application using the NodePort
- 6. Enabling autoscaling on a WordPress Pod to ensure scaling if the CPU utilization is greater than 50 percent